

Clustered Columnstore Indexes (SQL Server 2014)

Dave Fackler
Business Intelligence Architect
davef@rollinghillsky.com




Agenda

- What are (clustered) columnstore indexes?
- Some limitations
- How do inserts/updates/deletes work?
- Using clustered columnstore indexes

Columnstore Indexes

- Use new storage and compression based on columns (the Vertipaq engine)
 - Data split into row groups of 102,400 to 1,048,576 rows
 - Column segment created for each column in the table
 - Column segments (each column) compressed and stored



Types of Columnstore Indexes

- Non-clustered columnstore indexes
 - Introduced with SQL Server 2012
 - Can be built on table with clustered index or on a heap
 - Only includes those columns that are included when built
 - Makes the table read-only ☹️
- Clustered columnstore indexes
 - Introduced with SQL Server 2014
 - Is the only index on a table when created (no other indexes can be created)
 - Includes all columns of the table
 - Is updateable 😊

Limitations of Clustered Columnstore Indexes

- Cannot have any other indexes
- Table cannot have
 - PK constraints, FK constraints, unique constraints
 - Computed columns or sparse columns
 - Unsupported data types (text, uniqueidentifier, geography, xml, timestamp, hierarchyid, varchar(max), nvarchar(max))
- Cannot be combined with replication, change tracking, change data capture
- Cursors are not supported on CCI-indexed tables

Demo

(Scripts 1-12)

How do Inserts/Updates/Deletes Work?

- Clustered columnstore indexes use delta stores
 - Delta store is an “open” row group that can accept new rows
 - If delta store exists (less than 102,400 rows), it can be used
 - Otherwise, new delta stores are created as needed
- Inserts (unless using bulk insert) get added to delta store, but not immediately compressed
- Deletes get marked as deleted in columnstore segments
- Updates behave like deletes + inserts
- Delta stores get marked as “closed” once they are full

How do Inserts/Updates/Deletes Work?

- Closed delta stores compressed automatically over time
 - Tuple Mover background process looks for closed delta stores
 - Wakes up every few minutes, does work if needed, then sleeps
 - Is single-threaded, so it can take a long time to handle lots of closed delta stores...
- ALTER INDEX ... REORGANIZE
 - Works like Tuple Mover, but is multi-threaded
- ALTER INDEX ... REBUILD
 - Closes open delta stores (even if small)
 - Compresses all closed delta stores

Demo

(Scripts 13-18)

Using Clustered Columnstore Indexes

- Initially loading a large amount of data (100M+ rows)
 - Create table with regular clustered index, load, then create CCI
 - Or create table with CCI and bulk load data
- Monitor rowgroups and minimize open/closed delta stores
 - Queries must scan delta stores that are not compressed
 - Data there is not sorted, not indexed, and acts like a (slow) heap
 - Large number of rows in open/closed delta stores kills performance
- Use partitioning and reorganize/rebuild index by partition
 - After large sets of changes, determine partitions affected
 - Use REORGANIZE (faster) or REBUILD (more complete)

Using Clustered Columnstore Indexes

- Metadata information
 - `sys.column_store_segments`
 - Returns information about each column in each columnstore segment
 - `sys.column_store_dictionaries`
 - Returns information about dictionaries used (for columns needing them)
 - `sys.column_store_row_groups`
 - Returns information about rowgroups (likely the most useful of the three)
 - `sp_spaceused`
 - Warning: can return incorrect information for a table with CCI

Questions?