



SQL Server 2014


in-memory programming



Presented by: Miguel Cebollero @SQLMiguel






My Background

- Miguel E Cebollero
 -  @SQLMiguel
 -  in/SQLMiguel
 - SQLMiguel@Gmail.com
- Industry Experience:
 - Author, Speaker, Banking, Insurance, Sports, Software Co, Logistics and Legal Industries
- Speaker:
 - Local, Regional, National Conferences
- Community Volunteer:
 - Big Brothers / Big Sisters
 - Metropolitan Ministries
 - Habitat for Humanity



- Author: Pro T-SQL Programmers Guide; ApressMedia – Published: 02/26/2015
- 16 years in IT Specializing in Databases
 - Manager
 - Database Admin
 - Architect
 - Developer
- Education:
 -  Bachelor of Science, The University of Tampa
 -  Master of Science, The University of North Carolina Greensboro

Father and Husband



Presented by: Miguel Cebollero

In-Memory Development – SQL2014

1
Let's Get Started

2
Programming

3
Indexing

New TSQL development capabilities

Presented by: Miguel Cebollero @SQLMiguel




Background and Architecture

Hardware, Durability, Architecture

Presented by: Miguel Cebollero @SQLMiguel

What is it? How do I get it?



The diagram consists of three blue rounded rectangular boxes arranged horizontally. The first box on the left contains the text 'In-Memory OLTP'. The middle box contains 'In-Memory DW'. The third box on the right contains 'SSD Bufferpool Extension'. Each box has a subtle reflection effect below it.


SQL Server 2014 In-Memory features are offered in Enterprise, Developer or Evaluation (64-bit only)

Presented by: Miguel Cebollero @SQLMiguel

- In-Memory OLTP: The transactional lock-free, latch-free processing of transactional data on the database within the memory (RAM) of the server.
- In-Memory DW: Columnstore indexes that are contained within the memory (RAM) of the server.
- SSD Bufferpool Extension: provides seamless integration of the database engine buffer pool to improve I/O throughput.

Have a Time Machine? Grab your Laptop!

Year	Average Cost per Gigabyte
1980	\$6,635,520.00
1985	\$ 901,120.00
1990	\$ 108,544.00
1995	\$ 31,641.60
2000	\$ 1,149.95
2005	\$ 189.44
2010	\$ 12.50
2014	\$ 9.34



The cost of this laptop in 1980: \$53,084,160

Presented by: Miguel Cebollero @SQLMiguel

Why bother creating an in-memory database capability?

- Cost of Memory continues to plummet as the amount of memory continues to rise
- Microprocessor have hit a stagnant clock rate

- Source: Historical RAM Price; <http://www.statisticbrain.com/average-historic-price-of-ram/>

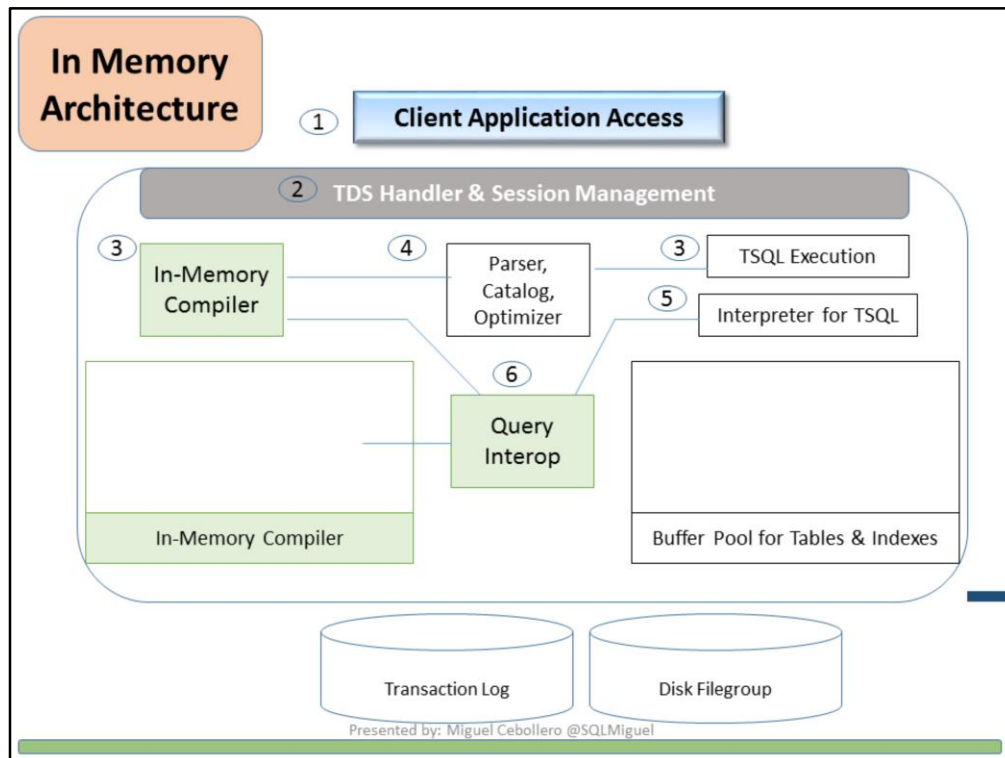
- Source: Moore's law; http://en.wikipedia.org/wiki/Moore's_law

Architecture

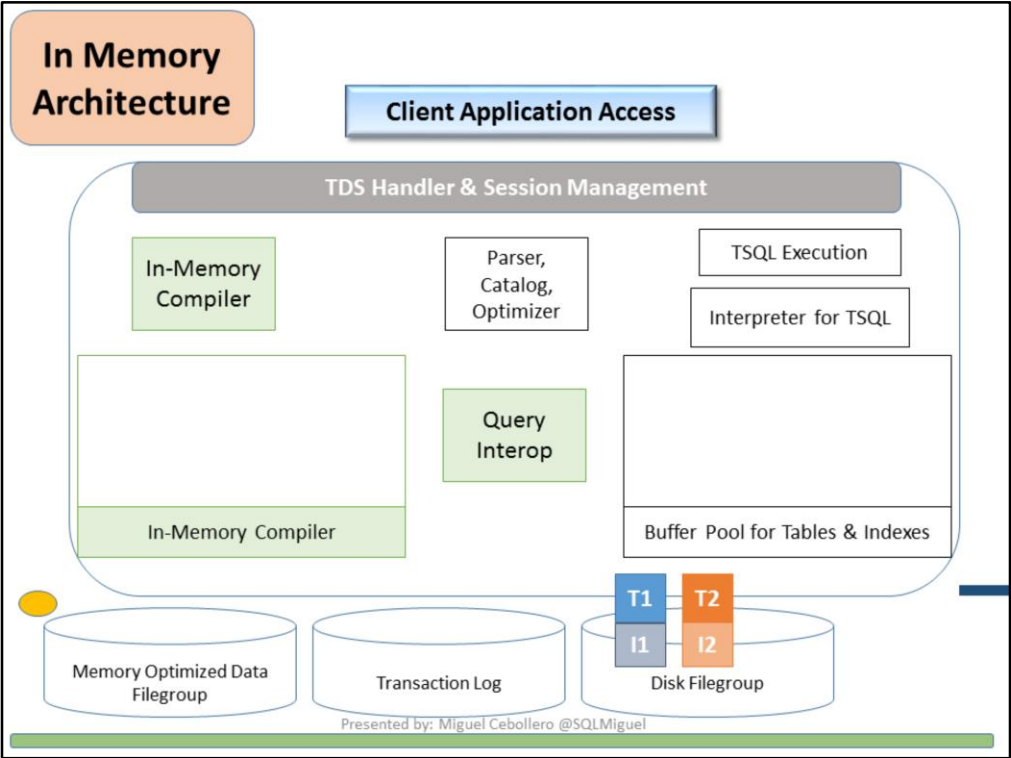
- **New row format**
 - No data pages
 - Rows are versioned, rather than being updated
 - Optimistic concurrency
 - Hard row size limit of 8060bytes
- **New Indexes**
 - Range Indexes
 - Hash Indexes
- **Compiled to DLL**

Presented by: Miguel Cebollero @SQLMiguel

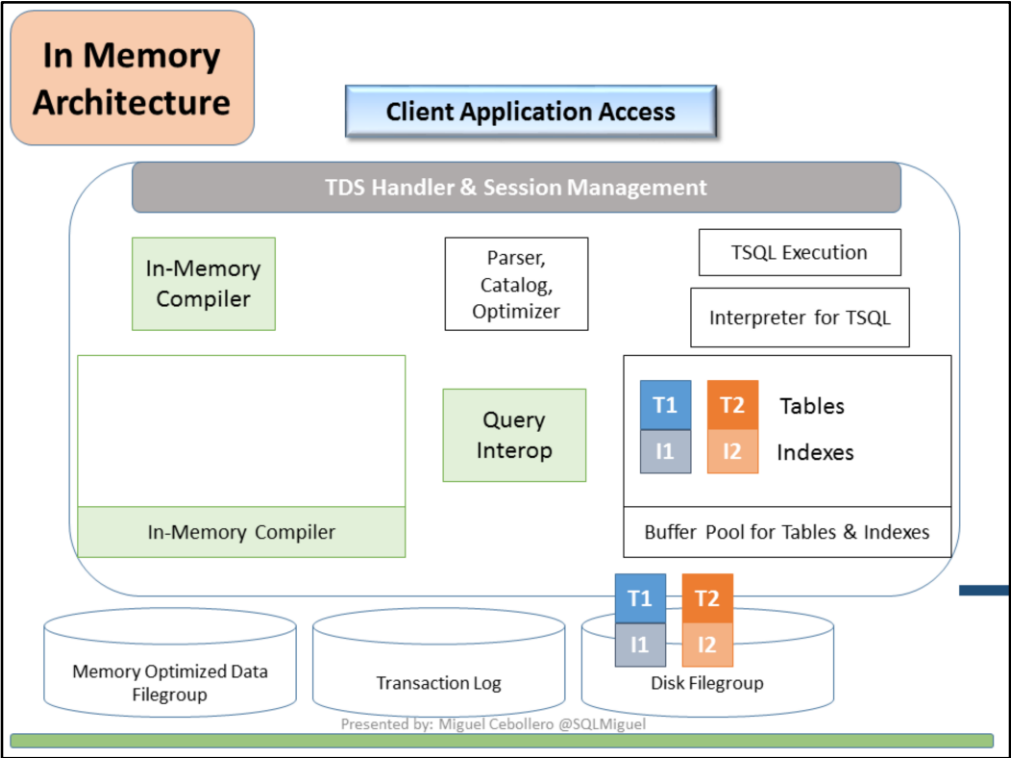
- Rows are linked via indexes in a linked list.



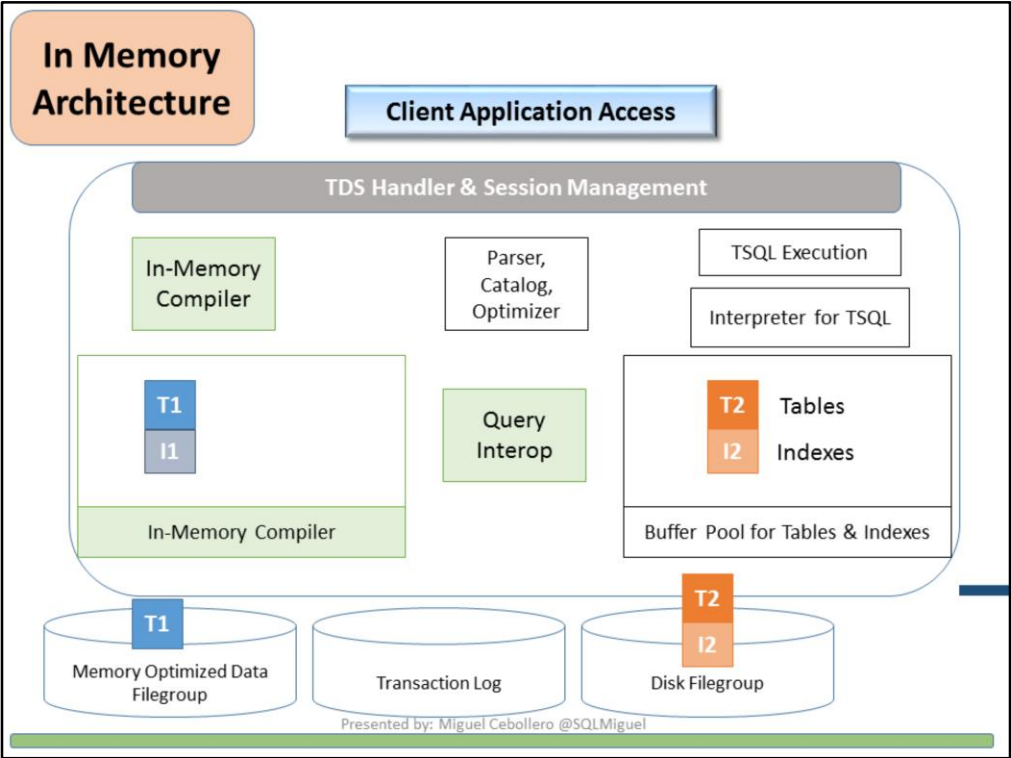
- 1. Client applications access SQL Server
 - 2. Through a single TDS Handler / Session Manager
 - 3. The handler decides to either access the In-Memory Compiler or TSQL Execution
 - 4. The Parser is used by either In-Memory or TSQL Execution
 - 5. Existing functionality
 - 6. Query Interop allows traditional SQL access to new In-Memory tables
- The optimizer is key, as it understands how to navigate new in-memory tables or stored procedures
- Source of Diagram: SQL Server 2014 Mission Critical Performance Level 300 Deck.pdf; Microsoft
 - Source: Microsoft SQL Server 2014 Hekaton CTP1 White Paper



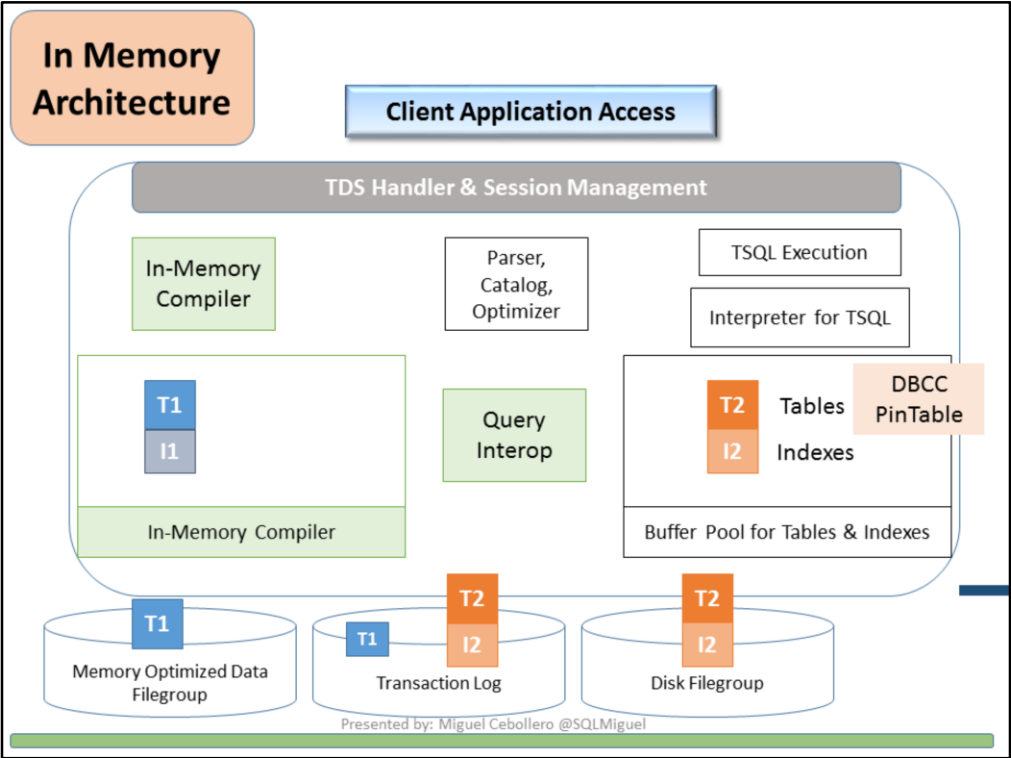
- Single Memory Optimized Data Filegroup is required
- Note: This filegroup can only be removed by dropping the database.




- Normal disk-based tables use the existing Buffer Pool to hold the rows of data needed to answer the current Client application request.



- Table-1 was re-created as an In-Memory table.
- Notice the indexes ONLY reside in memory; therefore, not in the Filegroup. More efficient than Disk-based for this reason.
- Only the index metadata / schema resides on disk. All of the table and index data is in memory.
- On restart the index and table data is moved off of disk back into memory.



- When writing to your tables there are differences between Disk-based and In-Memory
 - In-memory no index data being updated in the transaction log
 - Data written to Trans log is smaller
 - Data written to the filegroup is sequential and append only format
- Non-Durable tables will not generate any transaction logs and will not be written to the Filegroup.

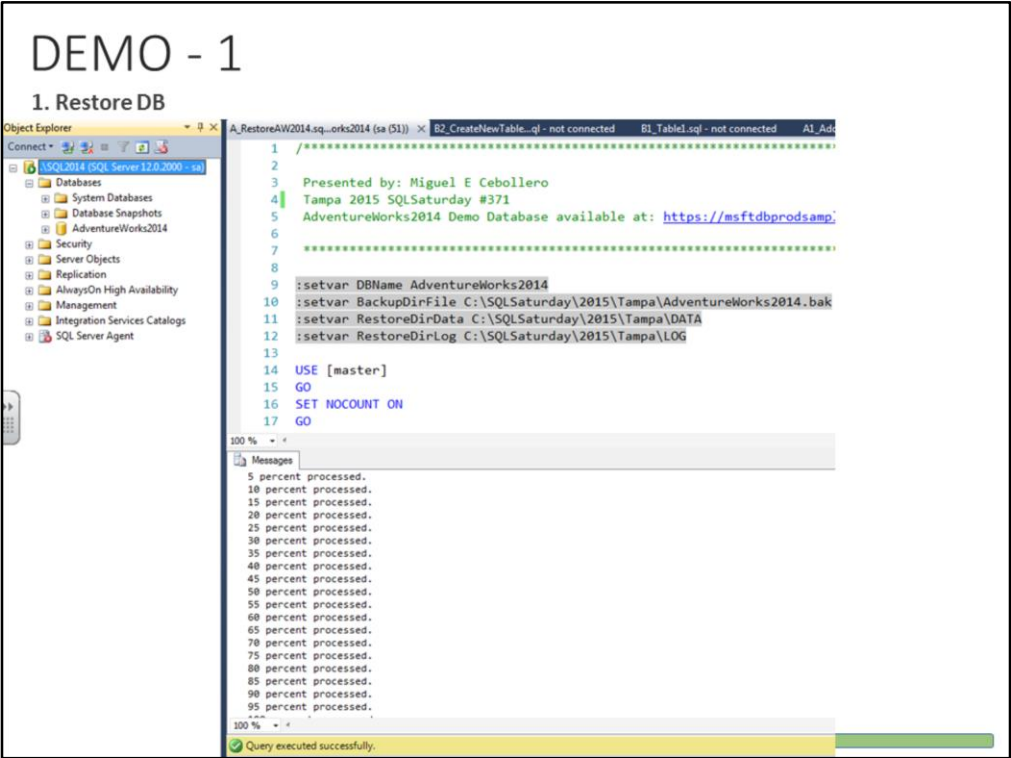


Let's Get Started

DEMO

Hardware, Durability, Architecture

Presented by: Miguel Cebollero @SQLMiguel



/

Presented by: Miguel E Cebollero
Tampa 2015 SQLSaturday #371
AdventureWorks2014 Demo Database available at:
https://msftdbprodsamples.codeplex.com/releases/view/125550

*****/

:setvar DBName AdventureWorks2014
:setvar BackupDirFile C:\SQLSaturday\2015\Tampa\AdventureWorks2014.bak
:setvar RestoreDirData C:\SQLSaturday\2015\Tampa\DATA
:setvar RestoreDirLog C:\SQLSaturday\2015\Tampa\LOG

USE [master]
GO
SET NOCOUNT ON
GO

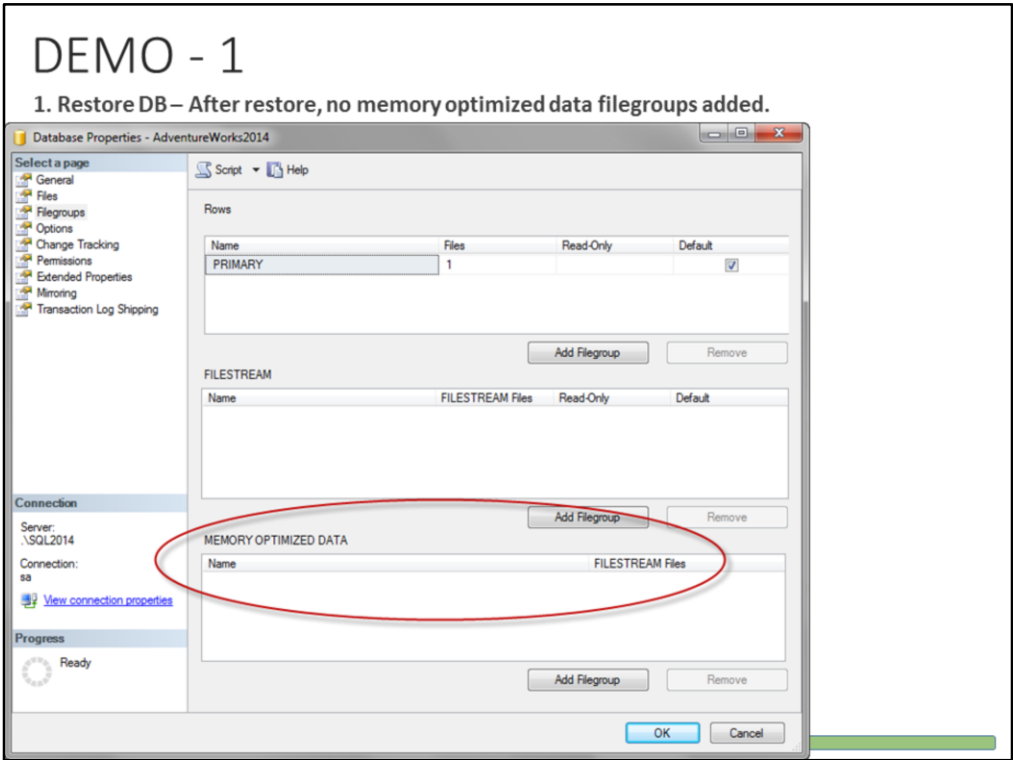
```
IF EXISTS (SELECT * FROM sys.databases where name='${DBName}')
DROP DATABASE [${DBName}];
GO
```

```
-- Restore AdventureWorks2014
RESTORE DATABASE [${DBName}]
FROM DISK = N'${BackupDirFile}'
WITH FILE = 1
, MOVE N'AdventureWorks2014_Data' TO
N'${RestoreDirData}\AdventureWorks2014_Data.mdf'
, MOVE N'AdventureWorks2014_Log' TO
N'${RestoreDirLog}\AdventureWorks2014_Log.ldf'
, NOUNLOAD, REPLACE, STATS = 5;
GO
```

```
ALTER AUTHORIZATION ON DATABASE::[${DBName}] TO [SA];
GO
```

```
ALTER DATABASE [${DBName}]
SET MULTI_USER
GO
```

```
USE AdventureWorks2014;
GO
UPDATE STATISTICS [Person].[Address];
GO
```



/*****
*****/

Presented by: Miguel E Cebollero
Tampa 2015 SQLSaturday #371
AdventureWorks2014 Demo Database available at:
<https://msftdbprodsamples.codeplex.com/releases/view/125550>

*****/

```

:setvar DBName AdventureWorks2014
:setvar BackupDirFile C:\SQLSaturday\2015\Tampa\AdventureWorks2014.bak
:setvar RestoreDirData C:\SQLSaturday\2015\Tampa\DATA
:setvar RestoreDirLog C:\SQLSaturday\2015\Tampa\LOG

```

```

USE [master]
GO
SET NOCOUNT ON
GO

```



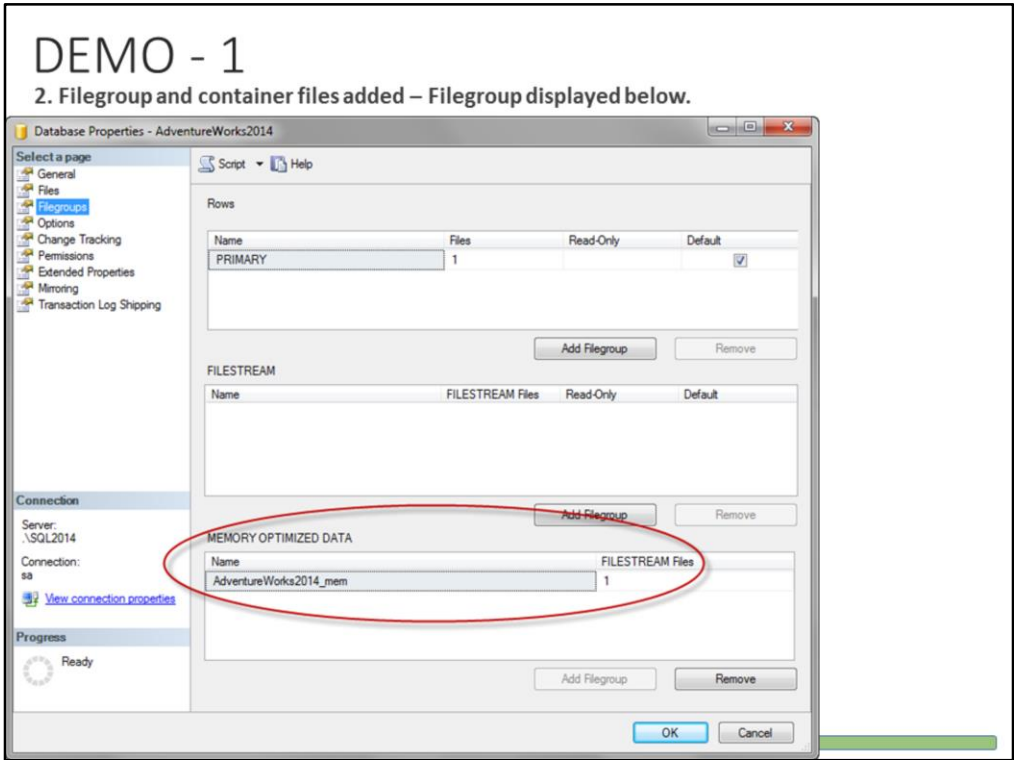
```
IF EXISTS (SELECT * FROM sys.databases where name='${DBName}')
DROP DATABASE [${DBName}];
GO
```

```
-- Restore AdventureWorks2014
RESTORE DATABASE [${DBName}]
FROM DISK = N'${BackupDirFile}'
WITH FILE = 1
, MOVE N'AdventureWorks2014_Data' TO
N'${RestoreDirData}\AdventureWorks2014_Data.mdf'
, MOVE N'AdventureWorks2014_Log' TO
N'${RestoreDirLog}\AdventureWorks2014_Log.ldf'
, NOUNLOAD, REPLACE, STATS = 5;
GO
```

```
ALTER AUTHORIZATION ON DATABASE::[${DBName}] TO [SA];
GO
```

```
ALTER DATABASE [${DBName}]
SET MULTI_USER
GO
```

```
USE AdventureWorks2014;
GO
UPDATE STATISTICS [Person].[Address];
GO
```



/*

Notes: Add MEMORY_OPTIMIZED_DATA filegroup and container to enable in-memory OLTP in the database
This code can be used to alter other existing databases.

*****/

```
:setvar DBName "AdventureWorks2014"  
:setvar CheckPointDir "C:\SQLSaturday\2015\Tampa\DATA\  
  
-- Create FILEGROUP  
-- Can only have one memory optimized filegroup per database  
IF NOT EXISTS (SELECT * FROM $(DBName).sys.data_spaces WHERE type='FX') --<  
InMemory abbreviation!  
ALTER DATABASE $(DBName)  
    ADD FILEGROUP [$(DBName)_mem] CONTAINS MEMORY_OPTIMIZED_DATA --<  
*** This is where all the magic happens ***
```

GO

-- Create New File and add it to the new Filegroup

-- You can have multiple containers (files) per memory optimized filegroup

-- Advantage of multiple containers is for recovery purposes

IF NOT EXISTS (SELECT * FROM \$(DBName).sys.data_spaces ds

JOIN \$(DBName).sys.database_files df ON

ds.data_space_id=df.data_space_id WHERE ds.type='FX') --< InMemory abbreviation

ALTER DATABASE \$(DBName)

ADD FILE (name='\$(DBName)_mem', filename='\$(CheckPointDir)\$(DBName)_mem')

TO FILEGROUP [\$(DBName)_mem]

GO

-- Navigate to file directory to see new files created

-- C:\SQLSaturday\2015\Tampa\DATA

-- Take a look at the properties of the database via SSMS

/*

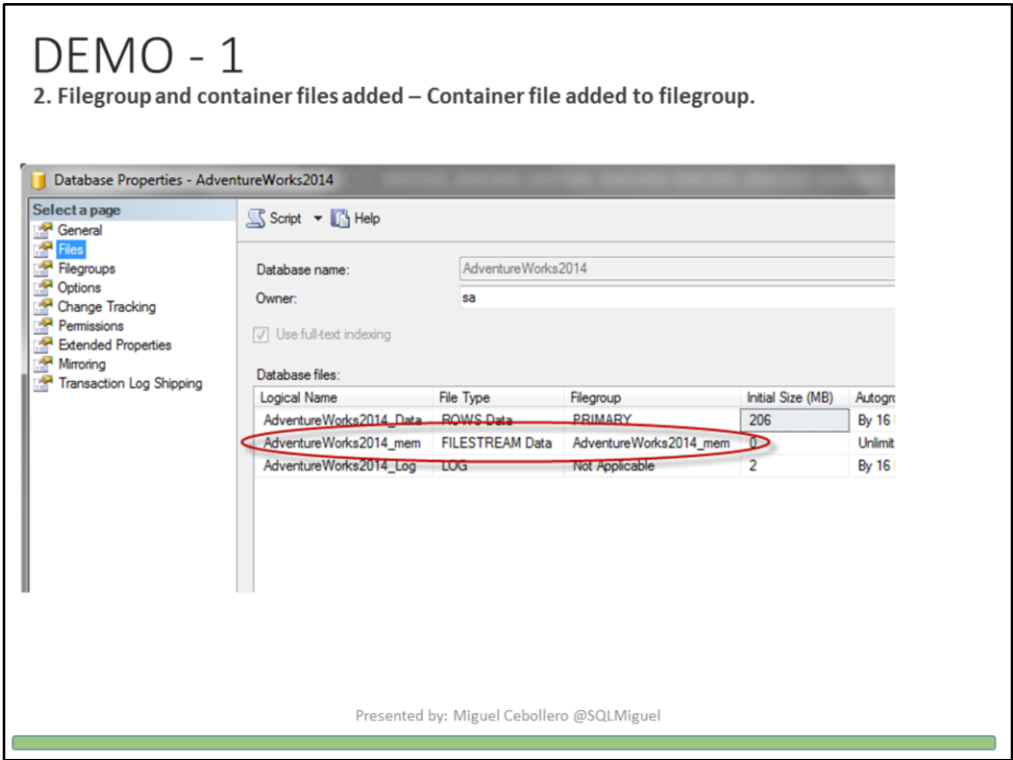
-- Take a look as the filegroup properties via system view

SELECT Name, Type, Type_Desc

FROM AdventureWorks2014.sys.data_spaces

WHERE type='FX' --< InMemory abbreviation

*/



/*

Notes: Add MEMORY_OPTIMIZED_DATA filegroup and container to enable in-memory OLTP in the database
This code can be used to alter other existing databases.

*****/

```
:setvar DBName "AdventureWorks2014"  
:setvar CheckPointDir "C:\SQLSaturday\2015\Tampa\DATA\  
  
-- Create FILEGROUP  
-- Can only have one memory optimized filegroup per database  
IF NOT EXISTS (SELECT * FROM $(DBName).sys.data_spaces WHERE type='FX') --<  
InMemory abbreviation!  
ALTER DATABASE $(DBName)  
    ADD FILEGROUP [$(DBName)_mem] CONTAINS MEMORY_OPTIMIZED_DATA --<  
*** This is where all the magic happens ***
```

GO

-- Create New File and add it to the new Filegroup

-- You can have multiple containers (files) per memory optimized filegroup

-- Advantage of multiple containers is for recovery purposes

IF NOT EXISTS (SELECT * FROM \$(DBName).sys.data_spaces ds

JOIN \$(DBName).sys.database_files df ON

ds.data_space_id=df.data_space_id WHERE ds.type='FX') --< InMemory abbreviation!

ALTER DATABASE \$(DBName)

ADD FILE (name='\$(DBName)_mem', filename='\$(CheckPointDir)\$(DBName)_mem')

TO FILEGROUP [\$(DBName)_mem]

GO

-- Navigate to file directory to see new files created

-- C:\SQLSaturday\2015\Tampa\DATA

-- Take a look at the properties of the database via SSMS

/*

-- Take a look as the filegroup properties via system view

SELECT Name, Type, Type_Desc

FROM AdventureWorks2014.sys.data_spaces

WHERE type='FX' --< InMemory abbreviation!

*/



Programming Programming

SQL Execution, Stored Procedures, In-Memory Tables

Presented by: Miguel Cebollero @SQLMiguel

In-Memory Tables

In-Memory Tables

- Rows are versioned
- Hard limitation of 8060bytes per row
- Only NonClustered Indexes
- Data not arranged in any specific manner
- No Data Pages; Link list pointers
- Native compilation to DLL
- Minimum of 1-Index is required

Durable

- Must have an indexed PRIMARY KEY
- DURABILITY = SCHEMA_AND_DATA
- Writes to Transaction Log
- Delta files written sequentially to disk

Non-Durable

- Must have an index
- DURABILITY = SCHEMA_ONLY
- No IO impact
- Completely in memory only, except for schema

Presented by: Miguel Cebollero @SQLMiguel

- Each row can have multiple versions, to allow concurrent reads and writes on the same row.
- The Indexes are what links the multiple rows into a table
- Tables cannot be altered. They would need to be dropped and recreated.
- <http://msdn.microsoft.com/en-us/library/dn511014.aspx>

In-Memory Table; Durable

```
CREATE TABLE [MOD].[Durable]
(
    TableID      INT NOT NULL PRIMARY KEY NONCLUSTERED
                -- InMemory tables can only have NONCLUSTERED Indexes
                -- By default SQL will try to create a PK as CLUSTERED
    , Column1    VARCHAR(24) NOT NULL
    , Column2    VARCHAR(24) NULL
) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
GO
```

This is the simplest form of an in-memory Durable table syntax.

Presented by: Miguel Cebollero @SQLMiguel

- Syntax, MSDN reference: <http://msdn.microsoft.com/en-us/library/dn133186.aspx>

In-Memory Table; NonDurable

```
CREATE TABLE [MOD].[NonDurable]
(
    TableID      INT NOT NULL
    -- New BIN2 Collation only necessary if in an index
    , Column1    VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
    INDEX [IX_Column1] ( [Column1] )
    , Column2    VARCHAR(24) NULL
) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
GO
```

This is the simplest form of an in-memory NonDurable table syntax.

Presented by: Miguel Cebollero @SQLMiguel

- Syntax, MSDN reference: <http://msdn.microsoft.com/en-us/library/dn133186.aspx>

In-Memory Stored Procedures

- Natively Compiled
- Can only access memory-optimized tables
- Not interpreted and produce DLLs
- Recompiled after database restart
- DLL not part of the database, just information to recreate them
- Parallel Processing
- Lives in SQL Server Memory

Presented by: Miguel Cebollero @SQLMiguel

- This offers 10x or more performance gains. More gains for more complicated scenarios.
- If you need to access a mix between disk-based and memory optimized tables, you need to use traditional interpreted stored procedures.

-

In-Memory Stored Procedures

```

-- Natively Compiled Stored Procedure Template
CREATE PROCEDURE <SP_Name>
(
    @Parameter1 DATETIME
    , @Parameter2 DATETIME )
WITH
    ① NATIVE_COMPILATION
    ② , SCHEMABINDING
    ③ , EXECUTE AS OWNER
AS
    ④ BEGIN ATOMIC
        WITH
    ⑤ ( TRANSACTION ISOLATION LEVEL = SNAPSHOT
        , LANGUAGE = 'english')
-- TSQL Goes Here

END;

```

1. Tells SQL Server this will be natively compiled
2. Ties execution to the schema of the object
3. Hardcodes execution rights
4. All or nothing succeeds in the procedure
5. It wants to hardcode these values at compile time.

Presented by: Miguel Cebollero @SQLMiguel

- This offers 10x or more performance gains. More gains for more complicated scenarios.
- If you need to access a mix between disk-based and memory optimized tables, you need to use traditional interpreted stored procedures.

-



DEMO

SQL Execution, Stored Procedures, In-Memory Tables

Presented by: Miguel Cebollero @SQLMiguel

DEMO - 2

B1. Create memory optimized Durable table – PK missing for a Durable table

```

36 *** Step-2 -- Create Memory Optimized Table
37
38 DROP TABLE [MOD].[Durable];
39 GO
40
41 CREATE TABLE [MOD].[Durable]
42 (
43     TableID INT NOT NULL
44     , Column1 VARCHAR(24) NOT NULL
45     , Column2 VARCHAR(24) NULL
46 )
47 --< 1. Note the lack of a filegroup
48 -- By default there is only one memory optimized filegroup; therefore, no need to specify
49 ) WITH(MEMORY_OPTIMIZED=ON); -- 2. , DURABILITY=SCHEMA_AND_DATA);
50     -- DURABILITY option NOT mandatory; default is SCHEMA_AND_DATA
51 GO
52 -- Received error, because you are forced to have a PK when creating a
53
54
100 %
Messages
Msg 41321, Level 16, State 7, Line 41
The memory optimized table 'Durable' with DURABILITY=SCHEMA_AND_DATA must have a primary key.
Msg 1750, Level 16, State 0, Line 41
Could not create constraint or index. See previous errors.

Presented by: Miguel Cebollero @SQLMiguel

```

USE AdventureWorks2014;
GO

/*****

Create new Memory Optimized Table

*****/

-- Create new schema for the purposes of comparison with disk-based table
CREATE SCHEMA [MOD] AUTHORIZATION [dbo]; -- Memory Optimized Data
GO

*** Step-1 -- Create Disk-based table

-- DROP TABLE [MOD].[Durable]

CREATE TABLE [MOD].[Durable]
(
TableID INT NOT NULL
, Column1 VARCHAR(24) NOT NULL

```
, Column2VARCHAR(24) NULL
, Column3VARCHAR(4000) NULL
, Column4VARCHAR(4000) NULL

) ON [PRIMARY] --< Note the filegroup
GO
-- No indexes, No PK, Just a Heap
```

*** Step-2 -- Create Memory Optimized Table

```
DROP TABLE [MOD].[Durable];
GO
```

```
CREATE TABLE [MOD].[Durable]
(
  TableIDINT NOT NULL
, Column1VARCHAR(24) NOT NULL
, Column2VARCHAR(24) NULL
```

```
--< 1. Note the lack of a filegroup
) WITH(MEMORY_OPTIMIZED=ON); -- 2. , DURABILITY=SCHEMA_AND_DATA);
-- DURABILITY option NOT mandatory
-- default is SCHEMA_AND_DATA
GO
```

*** Step-3 -- Missing PRIMARY KEY

```
CREATE TABLE [MOD].[Durable]
(
  TableID INT NOT NULL PRIMARY KEY --< Add our Primary Key
  , Column1 VARCHAR(24) NOT NULL
  , Column2 VARCHAR(24) NULL
) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
GO
```

*** Step-4 -- PRIMARY KEY fix; NONCLUSTERED

-- Simplest syntax for a DURABLE

--in-memory table.

```
CREATE TABLE [MOD].[Durable]
(
  TableID INT NOT NULL PRIMARY KEY NONCLUSTERED
  -- InMemory tables can only have NONCLUSTERED Indexes
  -- By default SQL will try to create a PK as CLUSTERED
  , Column1 VARCHAR(24) NOT NULL
  , Column2 VARCHAR(24) NULL
) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
```

GO

*** Step-5 -- NonDurable

```
CREATE TABLE [MOD].[NonDurable]
(
  TableID INT NOT NULL
-- 1) New Inline syntax for SQL2014
, Column1 VARCHAR(24) NOT NULL INDEX [IX_Column1] ( [Column1] )
, Column2 VARCHAR(24) NULL
, Column3 VARCHAR(4000) NULL
, Column4 VARCHAR(4000) NULL

, INDEX [IX_Column2] ( [Column2] ) -- 2) Can still declare index after columns
-- 3) NONCLUSTERED
--a. The "NONCLUSTERED" hint is optional,
--    unless being defined against the primary key

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
--< SCHEMA_ONLY option; literally no safety net
-- Use cases: Staging, Website Session State
GO
```


*** Step-6 -- NonDurable; Fix NULLABLE Column

```
CREATE TABLE [MOD].[NonDurable]
(
  TableID INT NOT NULL
, Column1 VARCHAR(24) NOT NULL INDEX [IX_Column1] ( [Column1] )
, Column2 VARCHAR(24) NOT NULL --< Fix NULLABLE Column
, Column3 VARCHAR(3850) NULL --< Fix our row size limitation issues
, Column4 VARCHAR(3850) NULL --< Fix our row size limitation issues

, INDEX [IX_Column2] ( [Column2] )

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
GO
```

*** Step-7 -- NonDurable; Fix BIN2 collation
-- Simplest form of a NonDurable table

```

CREATE TABLE [MOD].[NonDurable]
(
    TableID INT NOT NULL
-- New BIN2 Collation
, Column1 VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
INDEX [IX_Column1] ( [Column1] )
-- New BIN2 Collation only necessary if in an index
, Column2 VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
, Column3 VARCHAR(3850) NULL
, Column4 VARCHAR(3850) NULL

, INDEX [IX_Column2] ( [Column2] )
-- The string data type must be defined using a BIN2 collation

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
GO

```

```

/* -- Look at our new table in an existing system view
-- New 2014 columns are available

```

```

select t.name as 'Table Name'
, t.schema_id
, t.object_id
, filestream_data_space_id
, is_memory_optimized
, durability
, durability_desc
from sys.tables t
where type='U'
and t.schema_id = SCHEMA_ID(N'MOD');

```

```

*/

```

```

-- Look at our table via SSMS properties

```

```

-- Location of the dll files that represent the structure of the table we just created

```

```
-- C:\Program Files\Microsoft SQL Server\MSSQL12.SQL2014\MSSQL\DATA
```

DEMO - 2

B1. Create memory optimized Durable table – PK added, but clustered index not allowed

```

66 *** Step-3 -- Missing PRIMARY KEY
67
68 CREATE TABLE [MOD].[Durable]
69 (
70     TableID INT NOT NULL PRIMARY KEY --< Add our Primary Key
71     , Column1 VARCHAR(24) NOT NULL
72     , Column2 VARCHAR(24) NULL
73 )
74 WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
75 GO

```

100 %

Messages

Msg 12317, Level 16, State 72, Line 68
Clustered indexes, which are the default for primary keys, are not supported with memory optimized tables.

Presented by: Miguel Cebollero @SQLMiguel

USE AdventureWorks2014;
GO

/*****

Create new Memory Optimized Table

*****/

-- Create new schema for the purposes of comparison with disk-based table
CREATE SCHEMA [MOD] AUTHORIZATION [dbo]; -- Memory Optimized Data
GO

*** Step-1 -- Create Disk-based table

-- DROP TABLE [MOD].[Durable]

```

CREATE TABLE [MOD].[Durable]
(
    TableID INT NOT NULL
    , Column1 VARCHAR(24) NOT NULL

```

```
, Column2VARCHAR(24) NULL
, Column3VARCHAR(4000) NULL
, Column4VARCHAR(4000) NULL

) ON [PRIMARY] --< Note the filegroup
GO
-- No indexes, No PK, Just a Heap
```

*** Step-2 -- Create Memory Optimized Table

```
DROP TABLE [MOD].[Durable];
GO
```

```
CREATE TABLE [MOD].[Durable]
(
  TableIDINT NOT NULL
, Column1VARCHAR(24) NOT NULL
, Column2VARCHAR(24) NULL
```

```
--< 1. Note the lack of a filegroup
) WITH(MEMORY_OPTIMIZED=ON); -- 2. , DURABILITY=SCHEMA_AND_DATA);
-- DURABILITY option NOT mandatory
-- default is SCHEMA_AND_DATA
GO
```

*** Step-3 -- Missing PRIMARY KEY

```
CREATE TABLE [MOD].[Durable]
(
  TableID INT NOT NULL PRIMARY KEY --< Add our Primary Key
  , Column1 VARCHAR(24) NOT NULL
  , Column2 VARCHAR(24) NULL
) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
GO
```

*** Step-4 -- PRIMARY KEY fix; NONCLUSTERED

-- Simplest syntax for a DURABLE

--in-memory table.

```
CREATE TABLE [MOD].[Durable]
(
  TableID INT NOT NULL PRIMARY KEY NONCLUSTERED
  -- InMemory tables can only have NONCLUSTERED Indexes
  -- By default SQL will try to create a PK as CLUSTERED
  , Column1 VARCHAR(24) NOT NULL
  , Column2 VARCHAR(24) NULL
) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
```

GO

*** Step-5 -- NonDurable

```
CREATE TABLE [MOD].[NonDurable]
(
  TableID INT NOT NULL
-- 1) New Inline syntax for SQL2014
, Column1 VARCHAR(24) NOT NULL INDEX [IX_Column1] ( [Column1] )
, Column2 VARCHAR(24) NULL
, Column3 VARCHAR(4000) NULL
, Column4 VARCHAR(4000) NULL

, INDEX [IX_Column2] ( [Column2] ) -- 2) Can still declare index after columns
-- 3) NONCLUSTERED
--a. The "NONCLUSTERED" hint is optional,
--    unless being defined against the primary key

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
--< SCHEMA_ONLY option; literally no safety net
-- Use cases: Staging, Website Session State
GO
```

*** Step-6 -- NonDurable; Fix NULLABLE Column

```
CREATE TABLE [MOD].[NonDurable]
(
  TableID INT NOT NULL
, Column1 VARCHAR(24) NOT NULL INDEX [IX_Column1] ( [Column1] )
, Column2 VARCHAR(24) NOT NULL --< Fix NULLABLE Column
, Column3 VARCHAR(3850) NULL --< Fix our row size limitation issues
, Column4 VARCHAR(3850) NULL --< Fix our row size limitation issues

, INDEX [IX_Column2] ( [Column2] )

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
GO
```

*** Step-7 -- NonDurable; Fix BIN2 collation
-- Simplest form of a NonDurable table


```

CREATE TABLE [MOD].[NonDurable]
(
    TableID INT NOT NULL
-- New BIN2 Collation
, Column1 VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
INDEX [IX_Column1] ( [Column1] )
-- New BIN2 Collation only necessary if in an index
, Column2 VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
, Column3 VARCHAR(3850) NULL
, Column4 VARCHAR(3850) NULL

, INDEX [IX_Column2] ( [Column2] )
-- The string data type must be defined using a BIN2 collation

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
GO

```

```

/* -- Look at our new table in an existing system view
-- New 2014 columns are available

```

```

select t.name as 'Table Name'
, t.schema_id
, t.object_id
, filestream_data_space_id
, is_memory_optimized
, durability
, durability_desc
from sys.tables t
where type='U'
and t.schema_id = SCHEMA_ID(N'MOD');

```

```

*/

```

```

-- Look at our table via SSMS properties

```

```

-- Location of the dll files that represent the structure of the table we just created

```

```
-- C:\Program Files\Microsoft SQL Server\MSSQL12.SQL2014\MSSQL\DATA
```

DEMO - 2

B1. Create memory optimized Durable table – Successful; Simplest for of an in-memory table.

```

90
91 *** Step-4 -- PRIMARY KEY fix; NONCLUSTERED
92             -- Simplest syntax for a DURABLE
93             -- in-memory table.
94 CREATE TABLE [MOD].[Durable]
95 (
96     TableID    INT NOT NULL PRIMARY KEY NONCLUSTERED
97             -- InMemory tables can only have NONCLUSTERED Indexes
98             -- By default SQL will try to create a PK as CLUSTERED
99     , Column1  VARCHAR(24) NOT NULL
100     , Column2  VARCHAR(24) NULL
101 )
102 WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
103 GO
104
100 %
Messages
Command(s) completed successfully.

```

Presented by: Miguel Cebollero @SQLMiguel

USE AdventureWorks2014;
GO

/*****

Create new Memory Optimized Table

*****/

-- Create new schema for the purposes of comparison with disk-based table
CREATE SCHEMA [MOD] AUTHORIZATION [dbo]; -- Memory Optimized Data
GO

*** Step-1 -- Create Disk-based table

-- DROP TABLE [MOD].[Durable]

```

CREATE TABLE [MOD].[Durable]
(
    TableID INT NOT NULL
    , Column1 VARCHAR(24) NOT NULL

```

```
, Column2VARCHAR(24) NULL
, Column3VARCHAR(4000) NULL
, Column4VARCHAR(4000) NULL

) ON [PRIMARY] --< Note the filegroup
GO
-- No indexes, No PK, Just a Heap
```

*** Step-2 -- Create Memory Optimized Table

```
DROP TABLE [MOD].[Durable];
GO
```

```
CREATE TABLE [MOD].[Durable]
(
  TableIDINT NOT NULL
, Column1VARCHAR(24) NOT NULL
, Column2VARCHAR(24) NULL
```

```
--< 1. Note the lack of a filegroup
) WITH(MEMORY_OPTIMIZED=ON); -- 2. , DURABILITY=SCHEMA_AND_DATA);
-- DURABILITY option NOT mandatory
-- default is SCHEMA_AND_DATA
GO
```

*** Step-3 -- Missing PRIMARY KEY

```
CREATE TABLE [MOD].[Durable]
(
  TableID INT NOT NULL PRIMARY KEY --< Add our Primary Key
  , Column1 VARCHAR(24) NOT NULL
  , Column2 VARCHAR(24) NULL
) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
GO
```

*** Step-4 -- PRIMARY KEY fix; NONCLUSTERED

```
-- Simplest syntax for a DURABLE
--in-memory table.
CREATE TABLE [MOD].[Durable]
(
  TableID INT NOT NULL PRIMARY KEY NONCLUSTERED
  -- InMemory tables can only have NONCLUSTERED Indexes
  -- By default SQL will try to create a PK as CLUSTERED
  , Column1 VARCHAR(24) NOT NULL
  , Column2 VARCHAR(24) NULL
) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
```

GO

*** Step-5 -- NonDurable

```
CREATE TABLE [MOD].[NonDurable]
(
  TableID INT NOT NULL
-- 1) New Inline syntax for SQL2014
, Column1 VARCHAR(24) NOT NULL INDEX [IX_Column1] ( [Column1] )
, Column2 VARCHAR(24) NULL
, Column3 VARCHAR(4000) NULL
, Column4 VARCHAR(4000) NULL

, INDEX [IX_Column2] ( [Column2] ) -- 2) Can still declare index after columns
-- 3) NONCLUSTERED
--a. The "NONCLUSTERED" hint is optional,
--    unless being defined against the primary key

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
--< SCHEMA_ONLY option; literally no safety net
-- Use cases: Staging, Website Session State
GO
```

*** Step-6 -- NonDurable; Fix NULLABLE Column

```
CREATE TABLE [MOD].[NonDurable]
(
  TableID INT NOT NULL
, Column1 VARCHAR(24) NOT NULL INDEX [IX_Column1] ( [Column1] )
, Column2 VARCHAR(24) NOT NULL --< Fix NULLABLE Column
, Column3 VARCHAR(3850) NULL --< Fix our row size limitation issues
, Column4 VARCHAR(3850) NULL --< Fix our row size limitation issues

, INDEX [IX_Column2] ( [Column2] )

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
GO
```

*** Step-7 -- NonDurable; Fix BIN2 collation
-- Simplest form of a NonDurable table

```

CREATE TABLE [MOD].[NonDurable]
(
    TableID INT NOT NULL
-- New BIN2 Collation
, Column1 VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
INDEX [IX_Column1] ( [Column1] )
-- New BIN2 Collation only necessary if in an index
, Column2 VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
, Column3 VARCHAR(3850) NULL
, Column4 VARCHAR(3850) NULL

, INDEX [IX_Column2] ( [Column2] )
-- The string data type must be defined using a BIN2 collation

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
GO

```

```

/* -- Look at our new table in an existing system view
-- New 2014 columns are available

```

```

select t.name as 'Table Name'
, t.schema_id
, t.object_id
, filestream_data_space_id
, is_memory_optimized
, durability
, durability_desc
from sys.tables t
where type='U'
and t.schema_id = SCHEMA_ID(N'MOD');

```

```

*/

```

```

-- Look at our table via SSMS properties

```

```

-- Location of the dll files that represent the structure of the table we just created

```



```
-- C:\Program Files\Microsoft SQL Server\MSSQL12.SQL2014\MSSQL\DATA
```

DEMO - 2

B1. Create memory optimized NonDurable table – Cannot exceed a row size of 8060 bytes.

```

119 *** Step-5 -- NonDurable
120
121 CREATE TABLE [MOD].[NonDurable]
122 (
123     TableID INT NOT NULL
124     , Column1 VARCHAR(24) NOT NULL INDEX [IX_Column1] ( [Column1] )
125     , Column2 VARCHAR(24) NULL
126     , Column3 VARCHAR(5000) NULL
127     , Column4 VARCHAR(5000) NULL
128
129     , INDEX [IX_Column2] ( [Column2] ) -- 2) Can still declare index after columns
130     -- 3) NONCLUSTERED
131     -- a. The "NONCLUSTERED" hint is optional,
132     -- unless being defined against the primary key
133
134 ) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
135
136 --< SCHEMA_ONLY option; literally no safety net
137 -- Use cases: Staging, Website Session State
138 GO

```

100 %

Messages

Msg 41307, Level 16, State 1, Line 121
The row size limit of 8060 bytes for memory optimized tables has been exceeded. Please simplify the table definition.

Presented by: Miguel Cebollero @SQLMiguel

USE AdventureWorks2014;
GO

/*****

Create new Memory Optimized Table

*****/

-- Create new schema for the purposes of comparison with disk-based table
CREATE SCHEMA [MOD] AUTHORIZATION [dbo]; -- Memory Optimized Data
GO

*** Step-1 -- Create Disk-based table

-- DROP TABLE [MOD].[Durable]

```

CREATE TABLE [MOD].[Durable]
(
    TableID INT NOT NULL
    , Column1 VARCHAR(24) NOT NULL

```

```
, Column2VARCHAR(24) NULL
, Column3VARCHAR(4000) NULL
, Column4VARCHAR(4000) NULL

) ON [PRIMARY] --< Note the filegroup
GO
-- No indexes, No PK, Just a Heap
```

*** Step-2 -- Create Memory Optimized Table

```
DROP TABLE [MOD].[Durable];
GO
```

```
CREATE TABLE [MOD].[Durable]
(
  TableIDINT NOT NULL
, Column1VARCHAR(24) NOT NULL
, Column2VARCHAR(24) NULL
```

```
--< 1. Note the lack of a filegroup
) WITH(MEMORY_OPTIMIZED=ON); -- 2. , DURABILITY=SCHEMA_AND_DATA);
-- DURABILITY option NOT mandatory
-- default is SCHEMA_AND_DATA
GO
```

*** Step-3 -- Missing PRIMARY KEY

```
CREATE TABLE [MOD].[Durable]
(
  TableID INT NOT NULL PRIMARY KEY --< Add our Primary Key
  , Column1 VARCHAR(24) NOT NULL
  , Column2 VARCHAR(24) NULL
) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
GO
```

*** Step-4 -- PRIMARY KEY fix; NONCLUSTERED

-- Simplest syntax for a DURABLE

--in-memory table.

```
CREATE TABLE [MOD].[Durable]
(
  TableID INT NOT NULL PRIMARY KEY NONCLUSTERED
  -- InMemory tables can only have NONCLUSTERED Indexes
  -- By default SQL will try to create a PK as CLUSTERED
  , Column1 VARCHAR(24) NOT NULL
  , Column2 VARCHAR(24) NULL
) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
```

GO

*** Step-5 -- NonDurable

```
CREATE TABLE [MOD].[NonDurable]
(
  TableID INT NOT NULL
-- 1) New Inline syntax for SQL2014
, Column1 VARCHAR(24) NOT NULL INDEX [IX_Column1] ( [Column1] )
, Column2 VARCHAR(24) NULL
, Column3 VARCHAR(4000) NULL
, Column4 VARCHAR(4000) NULL

, INDEX [IX_Column2] ( [Column2] ) -- 2) Can still declare index after columns
-- 3) NONCLUSTERED
--a. The "NONCLUSTERED" hint is optional,
--    unless being defined against the primary key

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
--< SCHEMA_ONLY option; literally no safety net
-- Use cases: Staging, Website Session State
GO
```

*** Step-6 -- NonDurable; Fix NULLABLE Column

```
CREATE TABLE [MOD].[NonDurable]
(
  TableID INT NOT NULL
, Column1 VARCHAR(24) NOT NULL INDEX [IX_Column1] ( [Column1] )
, Column2 VARCHAR(24) NOT NULL --< Fix NULLABLE Column
, Column3 VARCHAR(3850) NULL --< Fix our row size limitation issues
, Column4 VARCHAR(3850) NULL --< Fix our row size limitation issues

, INDEX [IX_Column2] ( [Column2] )

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
GO
```

*** Step-7 -- NonDurable; Fix BIN2 collation
-- Simplest form of a NonDurable table

```

CREATE TABLE [MOD].[NonDurable]
(
    TableID INT NOT NULL
-- New BIN2 Collation
, Column1 VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
INDEX [IX_Column1] ( [Column1] )
-- New BIN2 Collation only necessary if in an index
, Column2 VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
, Column3 VARCHAR(3850) NULL
, Column4 VARCHAR(3850) NULL

, INDEX [IX_Column2] ( [Column2] )
-- The string data type must be defined using a BIN2 collation

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
GO

```

```

/* -- Look at our new table in an existing system view
-- New 2014 columns are available

```

```

select t.name as 'Table Name'
, t.schema_id
, t.object_id
, filestream_data_space_id
, is_memory_optimized
, durability
, durability_desc
from sys.tables t
where type='U'
and t.schema_id = SCHEMA_ID(N'MOD');

```

```

*/

```

```

-- Look at our table via SSMS properties

```

```

-- Location of the dll files that represent the structure of the table we just created

```

```
-- C:\Program Files\Microsoft SQL Server\MSSQL12.SQL2014\MSSQL\DATA
```


DEMO - 2

B1. Create memory optimized NonDurable table – Indexes on strings, must have _BIN2 collation.

```

154
155
156 *** Step-6 -- NonDurable; Fix NULLABLE Column
157
158 CREATE TABLE [MOD].[NonDurable]
159 (
160     TableID INT NOT NULL
161     , Column1 VARCHAR(24) NOT NULL INDEX [IX_Column1] ( [Column1] )
162     , Column2 VARCHAR(24) NOT NULL --< Fix NULLABLE Column
163     , Column3 VARCHAR(3850) NULL --< Fix our row size limitation issues
164     , Column4 VARCHAR(3850) NULL --< Fix our row size limitation issues
165
166     , INDEX [IX_Column2] ( [Column2] )
167
168 ) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
169 GO
170
171
172
173

```

100 %

Messages

Msg 12328, Level 16, State 102, Line 158
Indexes on character columns that do not use a *_BIN2 collation are not supported with indexes on memory optimized tables.
Msg 1750, Level 16, State 0, Line 158
Could not create constraint or index. See previous errors.

Presented by: Miguel Cebollero @SQLMiguel

USE AdventureWorks2014;
GO

/*****

Create new Memory Optimized Table

*****/

-- Create new schema for the purposes of comparison with disk-based table
CREATE SCHEMA [MOD] AUTHORIZATION [dbo]; -- Memory Optimized Data
GO

*** Step-1 -- Create Disk-based table

-- DROP TABLE [MOD].[Durable]

```

CREATE TABLE [MOD].[Durable]
(
    TableID INT NOT NULL
    , Column1 VARCHAR(24) NOT NULL

```

```
, Column2VARCHAR(24) NULL
, Column3VARCHAR(4000) NULL
, Column4VARCHAR(4000) NULL

) ON [PRIMARY] --< Note the filegroup
GO
-- No indexes, No PK, Just a Heap
```

*** Step-2 -- Create Memory Optimized Table

```
DROP TABLE [MOD].[Durable];
GO
```

```
CREATE TABLE [MOD].[Durable]
(
    TableID INT NOT NULL
, Column1 VARCHAR(24) NOT NULL
, Column2 VARCHAR(24) NULL
```

```
--< 1. Note the lack of a filegroup
) WITH(MEMORY_OPTIMIZED=ON); -- 2. , DURABILITY=SCHEMA_AND_DATA);
-- DURABILITY option NOT mandatory
-- default is SCHEMA_AND_DATA
GO
```

*** Step-3 -- Missing PRIMARY KEY

```
CREATE TABLE [MOD].[Durable]
(
  TableID INT NOT NULL PRIMARY KEY --< Add our Primary Key
  , Column1 VARCHAR(24) NOT NULL
  , Column2 VARCHAR(24) NULL
) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
GO
```

*** Step-4 -- PRIMARY KEY fix; NONCLUSTERED

-- Simplest syntax for a DURABLE

--in-memory table.

```
CREATE TABLE [MOD].[Durable]
(
  TableID INT NOT NULL PRIMARY KEY NONCLUSTERED
  -- InMemory tables can only have NONCLUSTERED Indexes
  -- By default SQL will try to create a PK as CLUSTERED
  , Column1 VARCHAR(24) NOT NULL
  , Column2 VARCHAR(24) NULL
) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
```

GO

*** Step-5 -- NonDurable

```
CREATE TABLE [MOD].[NonDurable]
(
  TableID INT NOT NULL
-- 1) New Inline syntax for SQL2014
, Column1 VARCHAR(24) NOT NULL INDEX [IX_Column1] ( [Column1] )
, Column2 VARCHAR(24) NULL
, Column3 VARCHAR(4000) NULL
, Column4 VARCHAR(4000) NULL

, INDEX [IX_Column2] ( [Column2] ) -- 2) Can still declare index after columns
-- 3) NONCLUSTERED
--a. The "NONCLUSTERED" hint is optional,
--    unless being defined against the primary key

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
--< SCHEMA_ONLY option; literally no safety net
-- Use cases: Staging, Website Session State
GO
```

*** Step-6 -- NonDurable; Fix NULLABLE Column

```
CREATE TABLE [MOD].[NonDurable]
(
  TableID INT NOT NULL
, Column1 VARCHAR(24) NOT NULL INDEX [IX_Column1] ( [Column1] )
, Column2 VARCHAR(24) NOT NULL --< Fix NULLABLE Column
, Column3 VARCHAR(3850) NULL --< Fix our row size limitation issues
, Column4 VARCHAR(3850) NULL --< Fix our row size limitation issues

, INDEX [IX_Column2] ( [Column2] )

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
GO
```

*** Step-7 -- NonDurable; Fix BIN2 collation
-- Simplest form of a NonDurable table

```

CREATE TABLE [MOD].[NonDurable]
(
    TableID INT NOT NULL
-- New BIN2 Collation
, Column1 VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
INDEX [IX_Column1] ( [Column1] )
-- New BIN2 Collation only necessary if in an index
, Column2 VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
, Column3 VARCHAR(3850) NULL
, Column4 VARCHAR(3850) NULL

, INDEX [IX_Column2] ( [Column2] )
-- The string data type must be defined using a BIN2 collation

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
GO

```

```

/* -- Look at our new table in an existing system view
-- New 2014 columns are available

```

```

select t.name as 'Table Name'
, t.schema_id
, t.object_id
, filestream_data_space_id
, is_memory_optimized
, durability
, durability_desc
from sys.tables t
where type='U'
and t.schema_id = SCHEMA_ID(N'MOD');

```

```

*/

```

```

-- Look at our table via SSMS properties

```

```

-- Location of the dll files that represent the structure of the table we just created

```

```
-- C:\Program Files\Microsoft SQL Server\MSSQL12.SQL2014\MSSQL\DATA
```

DEMO - 2

B1. Create memory optimized NonDurable table – Simplest form of a NonDurable table.

```

182 *** Step-7 -- NonDurable; Fix BIN2 collation
183         -- Simplest form of a NonDurable table
184
185 CREATE TABLE [MOD].[NonDurable]
186 (
187     TableID INT NOT NULL
188     -- New BIN2 Collation
189     , Column1 VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
190     INDEX [IX_Column1] ( [Column1] )
191     -- New BIN2 Collation only necessary if in an index
192     , Column2 VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
193     , Column3 VARCHAR(3850) NULL
194     , Column4 VARCHAR(3850) NULL
195
196     , INDEX [IX_Column2] ( [Column2] )
197     -- The string data type must be defined using a BIN2 collation
198
199 ) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
200 GO
201

```

100 %
 Messages
 Command(s) completed successfully.

Presented by: Miguel Cebollero @SQLMiguel

USE AdventureWorks2014;
 GO

/*****

Create new Memory Optimized Table

*****/

-- Create new schema for the purposes of comparison with disk-based table
 CREATE SCHEMA [MOD] AUTHORIZATION [dbo]; -- Memory Optimized Data
 GO

*** Step-1 -- Create Disk-based table

-- DROP TABLE [MOD].[Durable]

```

CREATE TABLE [MOD].[Durable]
(
    TableID INT NOT NULL
    , Column1 VARCHAR(24) NOT NULL

```



```
, Column2VARCHAR(24) NULL
, Column3VARCHAR(4000) NULL
, Column4VARCHAR(4000) NULL

) ON [PRIMARY] --< Note the filegroup
GO
-- No indexes, No PK, Just a Heap
```

*** Step-2 -- Create Memory Optimized Table

```
DROP TABLE [MOD].[Durable];
GO
```

```
CREATE TABLE [MOD].[Durable]
(
  TableIDINT NOT NULL
, Column1VARCHAR(24) NOT NULL
, Column2VARCHAR(24) NULL
```

```
--< 1. Note the lack of a filegroup
) WITH(MEMORY_OPTIMIZED=ON); -- 2. , DURABILITY=SCHEMA_AND_DATA);
-- DURABILITY option NOT mandatory
-- default is SCHEMA_AND_DATA
GO
```

*** Step-3 -- Missing PRIMARY KEY

```
CREATE TABLE [MOD].[Durable]
(
  TableID INT NOT NULL PRIMARY KEY --< Add our Primary Key
, Column1 VARCHAR(24) NOT NULL
, Column2 VARCHAR(24) NULL
) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
GO
```

*** Step-4 -- PRIMARY KEY fix; NONCLUSTERED

-- Simplest syntax for a DURABLE

--in-memory table.

```
CREATE TABLE [MOD].[Durable]
(
  TableID INT NOT NULL PRIMARY KEY NONCLUSTERED
-- InMemory tables can only have NONCLUSTERED Indexes
-- By default SQL will try to create a PK as CLUSTERED
, Column1 VARCHAR(24) NOT NULL
, Column2 VARCHAR(24) NULL
) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
```

GO

*** Step-5 -- NonDurable

```
CREATE TABLE [MOD].[NonDurable]
(
  TableID INT NOT NULL
-- 1) New Inline syntax for SQL2014
, Column1 VARCHAR(24) NOT NULL INDEX [IX_Column1] ( [Column1] )
, Column2 VARCHAR(24) NULL
, Column3 VARCHAR(4000) NULL
, Column4 VARCHAR(4000) NULL

, INDEX [IX_Column2] ( [Column2] ) -- 2) Can still declare index after columns
-- 3) NONCLUSTERED
--a. The "NONCLUSTERED" hint is optional,
--    unless being defined against the primary key

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
--< SCHEMA_ONLY option; literally no safety net
-- Use cases: Staging, Website Session State
GO
```

*** Step-6 -- NonDurable; Fix NULLABLE Column

```
CREATE TABLE [MOD].[NonDurable]
(
  TableID INT NOT NULL
, Column1 VARCHAR(24) NOT NULL INDEX [IX_Column1] ( [Column1] )
, Column2 VARCHAR(24) NOT NULL --< Fix NULLABLE Column
, Column3 VARCHAR(3850) NULL --< Fix our row size limitation issues
, Column4 VARCHAR(3850) NULL --< Fix our row size limitation issues

, INDEX [IX_Column2] ( [Column2] )

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
GO
```

*** Step-7 -- NonDurable; Fix BIN2 collation
-- Simplest form of a NonDurable table

```

CREATE TABLE [MOD].[NonDurable]
(
    TableID INT NOT NULL
    -- New BIN2 Collation
    , Column1 VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
    INDEX [IX_Column1] ( [Column1] )
    -- New BIN2 Collation only necessary if in an index
    , Column2 VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
    , Column3 VARCHAR(3850) NULL
    , Column4 VARCHAR(3850) NULL

    , INDEX [IX_Column2] ( [Column2] )
    -- The string data type must be defined using a BIN2 collation

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
GO

```

```

/* -- Look at our new table in an existing system view
    -- New 2014 columns are available

```

```

select t.name as 'Table Name'
    , t.schema_id
    , t.object_id
    , filestream_data_space_id
    , is_memory_optimized
    , durability
    , durability_desc
    from sys.tables t
    where type='U'
        and t.schema_id = SCHEMA_ID(N'MOD');

```

```

*/

```

```

-- Look at our table via SSMS properties

```

```

-- Location of the dll files that represent the structure of the table we just created

```

```
-- C:\Program Files\Microsoft SQL Server\MSSQL12.SQL2014\MSSQL\DATA
```

DEMO - 2

B1. Create memory optimized table – Query sys.tables to see new table descriptions.

```
204
205 /* -- Look at our new table in an existing system view
206     -- New 2014 columns are available
207
208 |select t.name as 'Table Name'
209     , t.schema_id
210     , t.object_id
211     , filestream_data_space_id
212     , is_memory_optimized
213     , durability
214     , durability_desc
215   from sys.tables t
216  where type='U'
217        and t.schema_id = SCHEMA_ID(N'MOD');
218
219 */
220
221
```

100 %

Results Messages

	Table Name	schema_id	object_id	filestream_data_space_id	is_memory_optimized	durability	durability_desc
1	Durable	10	663673412	NULL	1	0	SCHEMA_AND_DATA
2	NonDurable	10	695673526	NULL	1	1	SCHEMA_ONLY

Presented by: Miguel Cebollero @SQLMiguel

USE AdventureWorks2014;
GO

/******

Create new Memory Optimized Table

*****/

-- Create new schema for the purposes of comparision with disk-based table
CREATE SCHEMA [MOD] AUTHORIZATION [dbo]; -- Memory Optimized Data
GO

*** Step-1 -- Create Disk-based table

-- DROP TABLE [MOD].[Durable]

CREATE TABLE [MOD].[Durable]
(
TableID INT NOT NULL
, Column1 VARCHAR(24) NOT NULL

```
, Column2VARCHAR(24) NULL
, Column3VARCHAR(4000) NULL
, Column4VARCHAR(4000) NULL

) ON [PRIMARY] --< Note the filegroup
GO
-- No indexes, No PK, Just a Heap
```

*** Step-2 -- Create Memory Optimized Table

```
DROP TABLE [MOD].[Durable];
GO
```

```
CREATE TABLE [MOD].[Durable]
(
    TableID INT NOT NULL
, Column1 VARCHAR(24) NOT NULL
, Column2 VARCHAR(24) NULL
```

```
--< 1. Note the lack of a filegroup
) WITH(MEMORY_OPTIMIZED=ON); -- 2. , DURABILITY=SCHEMA_AND_DATA);
-- DURABILITY option NOT mandatory
-- default is SCHEMA_AND_DATA
GO
```


*** Step-3 -- Missing PRIMARY KEY

```
CREATE TABLE [MOD].[Durable]
(
  TableID INT NOT NULL PRIMARY KEY --< Add our Primary Key
, Column1 VARCHAR(24) NOT NULL
, Column2 VARCHAR(24) NULL
) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
GO
```

*** Step-4 -- PRIMARY KEY fix; NONCLUSTERED

-- Simplest syntax for a DURABLE

--in-memory table.

```
CREATE TABLE [MOD].[Durable]
(
  TableID INT NOT NULL PRIMARY KEY NONCLUSTERED
-- InMemory tables can only have NONCLUSTERED Indexes
-- By default SQL will try to create a PK as CLUSTERED
, Column1 VARCHAR(24) NOT NULL
, Column2 VARCHAR(24) NULL
) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_AND_DATA);
```

GO

*** Step-5 -- NonDurable

```
CREATE TABLE [MOD].[NonDurable]
(
  TableID INT NOT NULL
-- 1) New Inline syntax for SQL2014
, Column1 VARCHAR(24) NOT NULL INDEX [IX_Column1] ( [Column1] )
, Column2 VARCHAR(24) NULL
, Column3 VARCHAR(4000) NULL
, Column4 VARCHAR(4000) NULL

, INDEX [IX_Column2] ( [Column2] ) -- 2) Can still declare index after columns
-- 3) NONCLUSTERED
--a. The "NONCLUSTERED" hint is optional,
--    unless being defined against the primary key

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
--< SCHEMA_ONLY option; literally no safety net
-- Use cases: Staging, Website Session State
GO
```

*** Step-6 -- NonDurable; Fix NULLABLE Column

```
CREATE TABLE [MOD].[NonDurable]
(
  TableID INT NOT NULL
, Column1 VARCHAR(24) NOT NULL INDEX [IX_Column1] ( [Column1] )
, Column2 VARCHAR(24) NOT NULL --< Fix NULLABLE Column
, Column3 VARCHAR(3850) NULL --< Fix our row size limitation issues
, Column4 VARCHAR(3850) NULL --< Fix our row size limitation issues

, INDEX [IX_Column2] ( [Column2] )

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
GO
```

*** Step-7 -- NonDurable; Fix BIN2 collation
-- Simplest form of a NonDurable table

```

CREATE TABLE [MOD].[NonDurable]
(
    TableID INT NOT NULL
-- New BIN2 Collation
, Column1 VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
INDEX [IX_Column1] ( [Column1] )
-- New BIN2 Collation only necessary if in an index
, Column2 VARCHAR(24) COLLATE Latin1_General_100_BIN2 NOT NULL
, Column3 VARCHAR(3850) NULL
, Column4 VARCHAR(3850) NULL

, INDEX [IX_Column2] ( [Column2] )
-- The string data type must be defined using a BIN2 collation

) WITH(MEMORY_OPTIMIZED=ON, DURABILITY=SCHEMA_ONLY);
GO

```

```

/* -- Look at our new table in an existing system view
-- New 2014 columns are available

```

```

select t.name as 'Table Name'
, t.schema_id
, t.object_id
, filestream_data_space_id
, is_memory_optimized
, durability
, durability_desc
from sys.tables t
where type='U'
and t.schema_id = SCHEMA_ID(N'MOD');

```

```

*/

```

```

-- Look at our table via SSMS properties

```

```

-- Location of the dll files that represent the structure of the table we just created

```

```
-- C:\Program Files\Microsoft SQL Server\MSSQL12.SQL2014\MSSQL\DATA
```

DEMO - 2

B4. Demo
Differences – Disk
vs memory
optimized

```
12 /*****
13
14     Demonstrate Range Index vs. Disk-Based NonClustered Index Seek
15     ** Include execution plan with results **
16
17 *****/
18
19 SET STATISTICS IO ON
20 -- Single point lookup. Difference between disk-based vs memory optimized
21 -- Performance due to the Key Lookup to pull back all other column data
22 SELECT * FROM [Person].[Address] WHERE ModifiedDate = '2013-12-21';
23 -- All inmem indexes are covering; therefore, no Key Lookup
24 SELECT * FROM [MOD].[Address] WHERE ModifiedDate = '2013-12-21';
25
```

100 %

Results Spatial results Messages Execution plan

Query 1: Query cost (relative to the batch): 99%
SELECT * FROM [Person].[Address] WHERE [ModifiedDate]=@1

SELECT
Cost: 0 %

Nested Loops
(Inner Join)
Cost: 0 %

Index Seek (NonClustered)
[Address].[IX_Address_ModifiedDate]
Cost: 4 %

Key Lookup (Clustered)
[Address].[FK_Address_AddressID]
Cost: 96 %

Query 2: Query cost (relative to the batch): 1%
SELECT * FROM [MOD].[Address] WHERE [ModifiedDate]=@1

SELECT
Cost: 0 %

Index Seek (NonClustered)
[Address].[IX_MODAddress_ModifiedDa...
Cost: 100 %

USE AdventureWorks2014;
GO

-- Single point lookup using a Range Index
CHECKPOINT
GO
DBCC DROPCLEANBUFFERS
GO
DBCC FREEPROCCACHE
GO

/*

Demonstrate Range Index vs. Disk-Based NonClustered Index Seek
** Include execution plan with results **

*/

SET STATISTICS IO ON
-- Single point lookup. Difference between disk-based vs memory optimized
-- Performance due to the Key Lookup to pull back all other column data

```
SELECT * FROM [Person].[Address] WHERE ModifiedDate = '2013-12-21';  
-- All inmem indexes are covering; therefore, no Key Lookup  
SELECT * FROM [MOD].[Address] WHERE ModifiedDate = '2013-12-21';
```

SET STATISTICS IO ON

```
-- Performance almost equal, because the data is returned by the NC Index  
SELECT ModifiedDate FROM [Person].[Address] WHERE ModifiedDate = '2013-07-31';  
SELECT ModifiedDate FROM [MOD].[Address] WHERE ModifiedDate = '2013-07-31';
```

-- Perform a range query

```
SELECT * FROM [Person].[Address] WHERE ModifiedDate BETWEEN '2013-12-01'  
AND '2013-12-21';
```

-- As expected the memory optimized table performs better

```
SELECT * FROM [MOD].[Address] WHERE ModifiedDate BETWEEN '2013-12-01' AND  
'2013-12-21';
```

-- NonClustered Index in ASC order for our MOD table.

```
--, ModifiedDateDATETIME NOT NULL INDEX [IX_MODAddress_ModifiedDate]  
NONCLUSTERED
```

-- Performs Index Seek

```
SELECT * FROM [MOD].[Address] WHERE ModifiedDate BETWEEN '2013-12-01' AND  
'2013-12-21' ORDER BY ModifiedDate;
```

-- Performs Index Seek, but must also sort the data; cannot return the data in
opposite order of the index.

```
SELECT * FROM [MOD].[Address] WHERE ModifiedDate BETWEEN '2013-12-01' AND  
'2013-12-21' ORDER BY ModifiedDate DESC; -- ASC;
```

DEMO - 2

B4. Demo Differences – Difference almost equal, due to covering index on disk-based table.

```
26
27 SET STATISTICS IO ON
28 -- Performance almost equal, because the data is returned by the NC Index on the
29 -- disk-based table. In-memory is not always faster
30 SELECT ModifiedDate FROM [Person].[Address] WHERE ModifiedDate = '2013-07-31';
31 SELECT ModifiedDate FROM [MOD].[Address] WHERE ModifiedDate = '2013-07-31';
32
```

100 %

ResultsMessagesExecution plan

Query 1: Query cost (relative to the batch): 56%

SELECT [ModifiedDate] FROM [Person].[Address] WHERE [ModifiedDate]=@1

SELECT

Cost: 0 %

Index Seek (NonClustered)

[Address].[IX_Address_ModifiedDate]

Cost: 100 %

Query 2: Query cost (relative to the batch): 44%

SELECT [ModifiedDate] FROM [MOD].[Address] WHERE [ModifiedDate]=@1

SELECT

Cost: 0 %

Index Seek (NonClustered)

[Address].[IX_MODAddress_ModifiedDa...

Cost: 100 %

USE AdventureWorks2014;
GO

-- Single point lookup using a Range Index
CHECKPOINT
GO
DBCC DROPCLEANBUFFERS
GO
DBCC FREEPROCCACHE
GO

/*

Demonstrate Range Index vs. Disk-Based NonClustered Index Seek
** Include execution plan with results **

*/

SET STATISTICS IO ON
-- Single point lookup. Difference between disk-based vs memory optimized
-- Performance due to the Key Lookup to pull back all other column data


```
SELECT * FROM [Person].[Address] WHERE ModifiedDate = '2013-12-21';
```

-- All inmem indexes are covering; therefore, no Key Lookup

```
SELECT * FROM [MOD].[Address] WHERE ModifiedDate = '2013-12-21';
```

```
SET STATISTICS IO ON
```

-- Performance almost equal, because the data is returned by the NC Index

```
SELECT ModifiedDate FROM [Person].[Address] WHERE ModifiedDate = '2013-07-31';
```

```
SELECT ModifiedDate FROM [MOD].[Address] WHERE ModifiedDate = '2013-07-31';
```

-- Perform a range query

```
SELECT * FROM [Person].[Address] WHERE ModifiedDate BETWEEN '2013-12-01'  
AND '2013-12-21';
```

-- As expected the memory optimized table performs better

```
SELECT * FROM [MOD].[Address] WHERE ModifiedDate BETWEEN '2013-12-01' AND  
'2013-12-21';
```

-- NonClustered Index in ASC order for our MOD table.

```
--, ModifiedDateDATETIME NOT NULL INDEX [IX_MODAddress_ModifiedDate]  
NONCLUSTERED
```

-- Performs Index Seek

```
SELECT * FROM [MOD].[Address] WHERE ModifiedDate BETWEEN '2013-12-01' AND  
'2013-12-21' ORDER BY ModifiedDate;
```

-- Performs Index Seek, but must also sort the data; cannot return the data in
opposite order of the index.

```
SELECT * FROM [MOD].[Address] WHERE ModifiedDate BETWEEN '2013-12-01' AND  
'2013-12-21' ORDER BY ModifiedDate DESC; -- ASC;
```

DEMO - 2

B4. Demo Differences – Difference almost equal, due to covering index on disk-based table.

```
26
27 SET STATISTICS IO ON
28 -- Performance almost equal, because the data is returned by the NC Index on the
29 -- disk-based table. In-memory is not always faster
30 SELECT ModifiedDate FROM [Person].[Address] WHERE ModifiedDate = '2013-07-31';
31 SELECT ModifiedDate FROM [MOD].[Address] WHERE ModifiedDate = '2013-07-31';
32
```

100 %

ResultsMessagesExecution plan

Query 1: Query cost (relative to the batch): 56%

SELECT [ModifiedDate] FROM [Person].[Address] WHERE [ModifiedDate]=@1

SELECT

Cost: 0 %

Index Seek (NonClustered)

[Address].[IX_Address_ModifiedDate]

Cost: 100 %

Query 2: Query cost (relative to the batch): 44%

SELECT [ModifiedDate] FROM [MOD].[Address] WHERE [ModifiedDate]=@1

SELECT

Cost: 0 %

Index Seek (NonClustered)

[Address].[IX_MODAddress_ModifiedDa...

Cost: 100 %

USE AdventureWorks2014;
GO

-- Single point lookup using a Range Index
CHECKPOINT
GO
DBCC DROPCLEANBUFFERS
GO
DBCC FREEPROCCACHE
GO

/*

Demonstrate Range Index vs. Disk-Based NonClustered Index Seek
** Include execution plan with results **

*/

SET STATISTICS IO ON
-- Single point lookup. Difference between disk-based vs memory optimized
-- Performance due to the Key Lookup to pull back all other column data

```
SELECT * FROM [Person].[Address] WHERE ModifiedDate = '2013-12-21';  
-- All inmem indexes are covering; therefore, no Key Lookup  
SELECT * FROM [MOD].[Address] WHERE ModifiedDate = '2013-12-21';
```

SET STATISTICS IO ON

```
-- Performance almost equal, because the data is returned by the NC Index  
SELECT ModifiedDate FROM [Person].[Address] WHERE ModifiedDate = '2013-07-31';  
SELECT ModifiedDate FROM [MOD].[Address] WHERE ModifiedDate = '2013-07-31';
```

-- Perform a range query

```
SELECT * FROM [Person].[Address] WHERE ModifiedDate BETWEEN '2013-12-01'  
AND '2013-12-21';
```

-- As expected the memory optimized table performs better

```
SELECT * FROM [MOD].[Address] WHERE ModifiedDate BETWEEN '2013-12-01' AND  
'2013-12-21';
```

-- NonClustered Index in ASC order for our MOD table.

```
--, ModifiedDateDATETIME NOT NULL INDEX [IX_MODAddress_ModifiedDate]  
NONCLUSTERED
```

-- Performs Index Seek

```
SELECT * FROM [MOD].[Address] WHERE ModifiedDate BETWEEN '2013-12-01' AND  
'2013-12-21' ORDER BY ModifiedDate;
```

-- Performs Index Seek, but must also sort the data; cannot return the data in
opposite order of the index.

```
SELECT * FROM [MOD].[Address] WHERE ModifiedDate BETWEEN '2013-12-01' AND  
'2013-12-21' ORDER BY ModifiedDate DESC; -- ASC;
```



Make Your Tables Scream

Hash Index, Range Index, Non-clustered

Presented by: Miguel Cebollero @SQLMiguel

In-Memory Table Indexes

Where the rubber
meets the road!

- Hash Index
 - Single Item lookups
 - Cannot be used in a LIKE operator
 - Bucket size determined at creating
- Range Index
 - Not as good for single item lookups
 - Good for range queries
 - Size of index grows with size of data

Presented by: Miguel Cebollero @SQLMiguel

Hash Index vs. Range Index

```
SELECT *  
FROM MOD.Address  
WHERE AddressId = 26007;
```

Hash Index

```
SELECT *  
FROM MOD.Address  
WHERE AddressId BETWEEN 1 AND 26007;
```

Range Index

Presented by: Miguel Cebollero @SQLMiguel

- This offers 10x or more performance gains. More gains for more complicated scenarios.
- If you need to access a mix between disk-based and memory optimized tables, you need to use traditional interpreted stored procedures.

-

Index Guidelines

- Minimum of 1-Index Per Table
- All Indexes are Non-Clustered
- Not Stored on Disk; Recreated during system recovery
- Cannot be Altered or Added after the table is created
- Cannot have a Unique Constraint; Only Primary Key
- Don't Duplicate Data, point to rows in a chain

Presented by: Miguel Cebollero @SQLMiguel

- This offers 10x or more performance gains. More gains for more complicated scenarios.
- If you need to access a mix between disk-based and memory optimized tables, you need to use traditional interpreted stored procedures.

-



DEMO

Hash Index, Range Index, Non-clustered

Presented by: Miguel Cebollero @SQLMiguel

References / Links

- <http://www.BIDataPartners.com>

