

Database Design Disasters

Richie Rump

Jorriss LLC

@Jorriss

www.jorriss.net

Who is this dude?



dotnetmiami

StatisticsParser.com

Why Database Design?

- Think of your database as your application foundation.
- The more time you spend on your foundation the better your application can be.
- If you get the foundation wrong you **WILL** have problems.

Why Database Design?

- Databases are highly complex systems.
- SQL Server has over twenty years of development.
- Learning SQL Server is like Alice's rabbit hole. It keeps going down.
- Leverage your DBAs experience and knowledge of SQL Server.

Why Database Design?

- Relational Databases are NOT going away!
- Reporting from NoSQL databases is difficult...for now.
- Right tool, right situation.

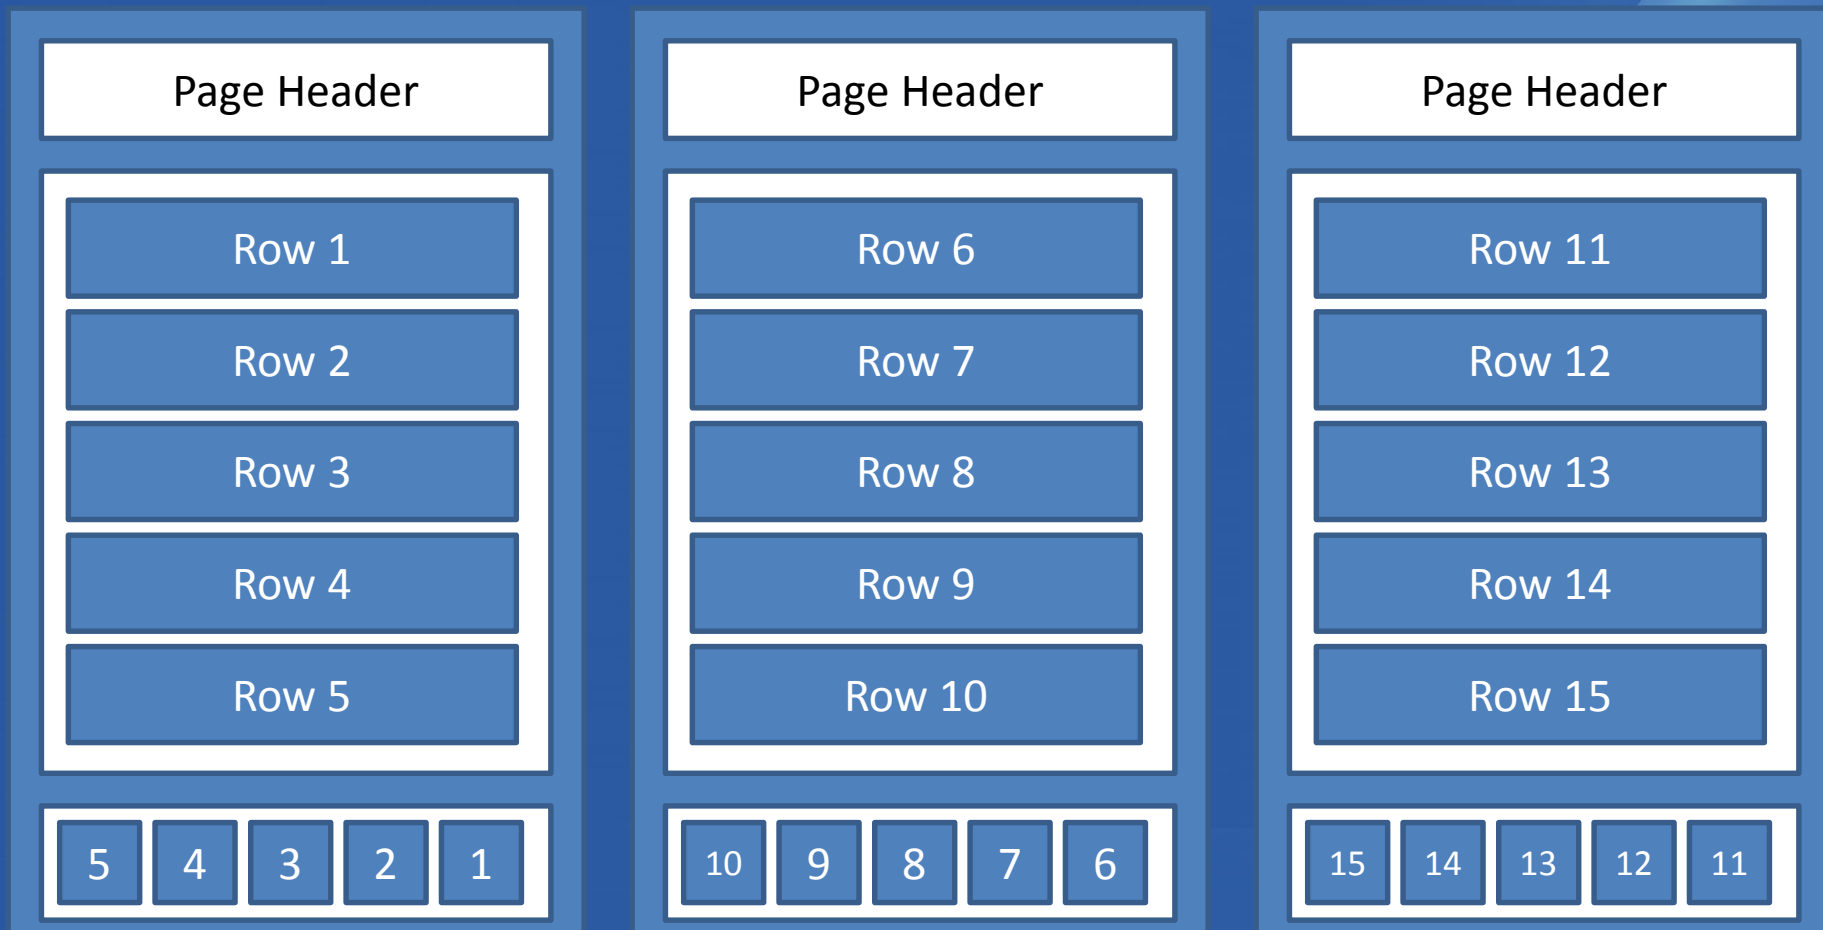
Current Toolset

- ORM tools like Hibernate, ActiveRecord and Entity Framework all can create databases.
- Out of the box schemas can have...opportunities.
- Agile frameworks and rapid iterations can leave DB design an afterthought.

Overview

- Wrong Types
- No Indexes on Foreign Keys
- Entity Attribute Value pattern
- Guids
- Surrogate Key / No Alternate Key

How SQL Server Stores Data



Poor Data Types

- Having the wrong data types can create long lasting issues.
- T-SQL can become challenging and inefficient.
- Understanding the differences between data types
- Better information, better decisions.

CHAR vs VARCHAR

- CHAR
 - Fixed length. Value will be padded with spaces.
 - Size *n* bytes.
 - 1 – 8000 characters.
- VARCHAR
 - Variable length.
 - The storage size is the actual length of the data entered + 2 bytes.
 - 1 – 8000 characters.

NCHAR vs NVARCHAR

- Holds UNICODE data.
- NCHAR
 - 1 – 4000 characters
 - Storage size twice n characters.
- NVARCHAR
 - 1 – 4000 character
 - Storage size twice n characters + 2 bytes.

NVARCHAR(MAX) / VARCHAR(MAX)

- MAX allows a column size of 2GB.
- Cannot create an index on MAX columns
- REPEAT: Cannot create an index on MAX columns
- NVARCHAR(MAX) is the default for a string when generated from Entity Framework.

DATETIME vs DATETIME2

- DATETIME
 - Accuracy rounded to .000, .003, or .007 second.
 - 8 Bytes
- DATETIME2 (*n*)
 - Accuracy to .0000001 second
 - 6 bytes for precisions less than 3;
 - 7 bytes for precisions 3 and 4.
 - All other precisions require 8 bytes.
- Demo

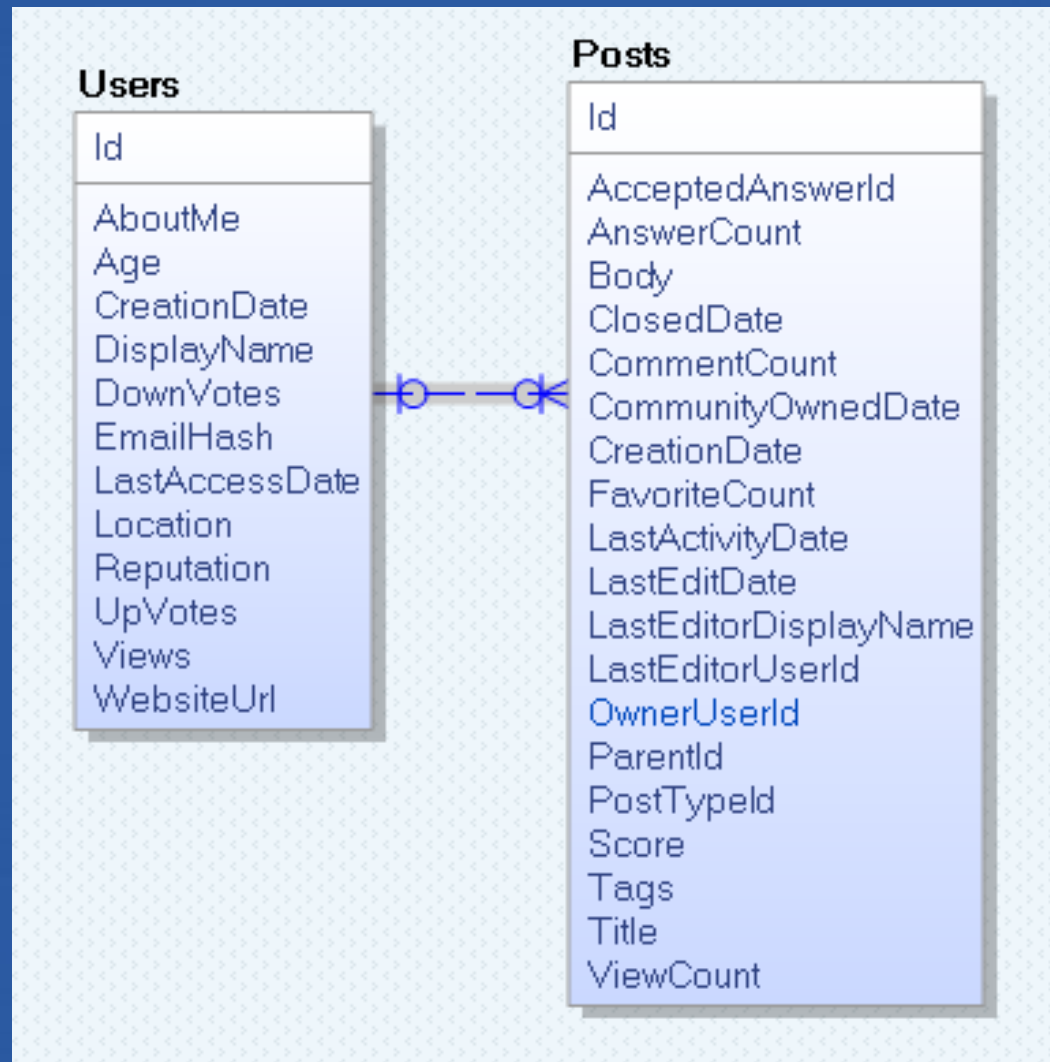
BIGINT vs INT

- BIGINT
 - -9,223,372,036,854,775,808 to 9,223,372,036,854,775,807
 - Size 8 bytes
- INT
 - -2,147,483,648 to 2,147,483,647
 - Size 4 bytes

No Indexes on Foreign Keys

- May slow queries with JOINS
- Will have a performance impact on DELETES

No Indexes on Foreign Keys




No Indexes on Foreign Keys





Demo

Entity-Attribute-Value



Entity

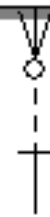
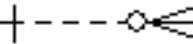
 Entity_ID INTEGER IDENTITY

Attribute

 Attribute_Id INTEGER IDENTITY	
 Entity_ID (FK) INTEGER	NOT NULL
 Attribute_Type_ID (FK) INTEGER	NOT NULL
 Value VARCHAR(max)	NULL

Attribute_Type

 Attribute_Type_ID INTEGER IDENTITY	
 Attribute_Name VARCHAR(50)	NULL



Entity-Attribute-Value

Person

Entity ID	Name
1	Cecil Phillip

Person_Attribute

Attribute ID	Entity ID	Attribute Type	Value
100	1	Phone Number	305-555-9607
101	1	Age	30
102	1	Birthdate	2/9/1983

Entity-Attribute-Value

- No Data Type Enforcement

Person

Entity ID	Name
1	Cecil Phillip

Person_Attribute

Attribute ID	Entity ID	Attribute Type	Value
100	1	Phone Number	305-555-9607
101	1	Age	501
102	1	Birthdate	Yes I have one

Entity-Attribute-Value

- No NOT NULL Enforcement

Person

Entity ID	Name
1	Cecil Phillip

Person_Attribute

Attribute ID	Entity ID	Attribute Type	Value
100	1	Phone Number	NULL
101	1	Age	30
102	1	Birthdate	2/9/1983

Entity-Attribute-Value

- Querying much harder
- Demo...

GUID Primary Keys

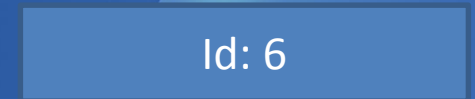
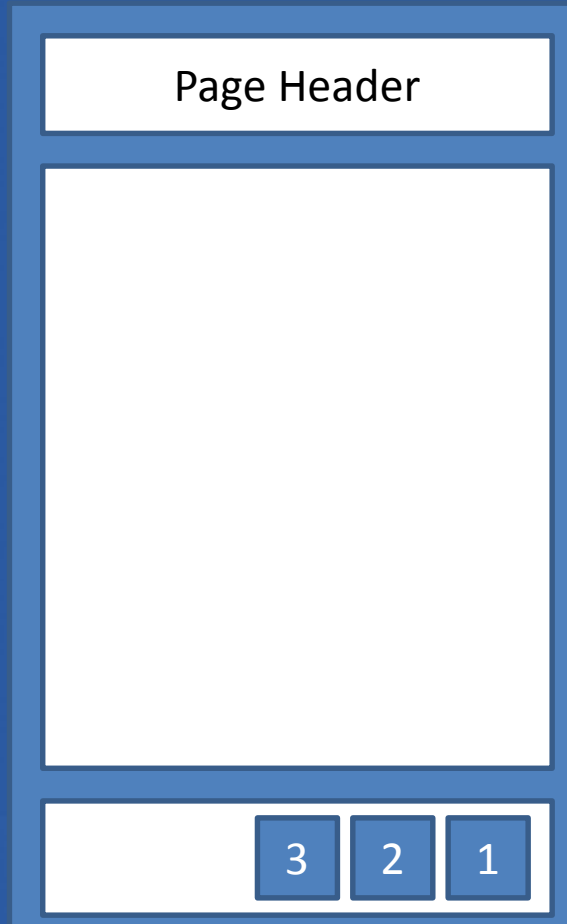
- You would think this wouldn't be a problem...
- Some benefits
 - Portable
 - Id can be generated from the app
 - That's about it.

GUID Primary Keys

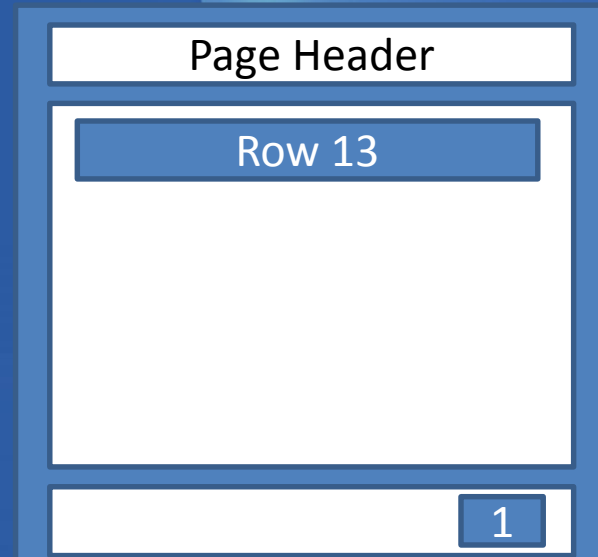
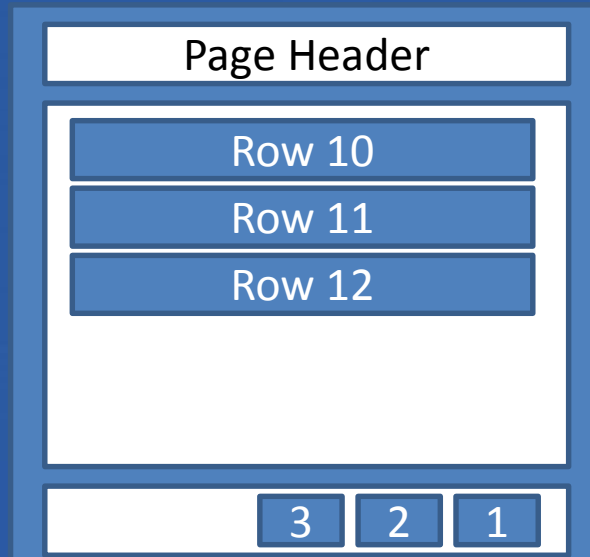
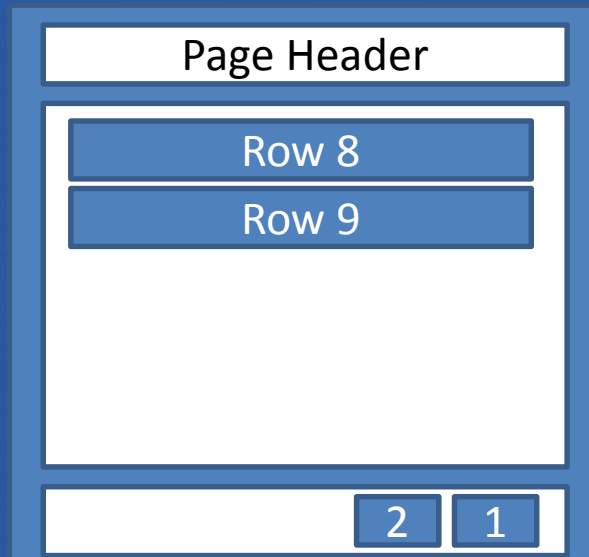
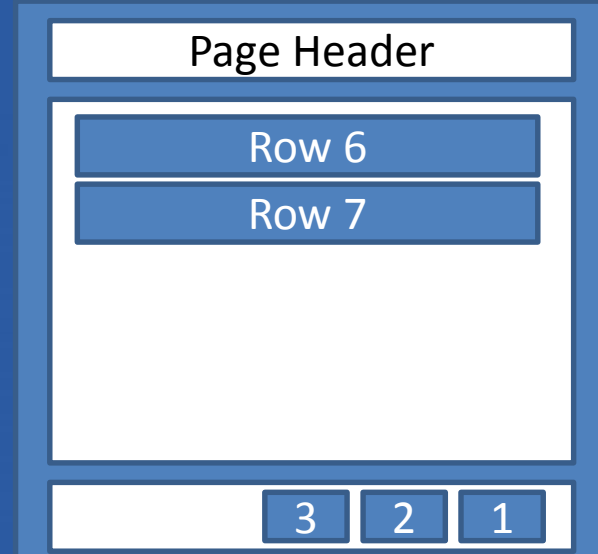
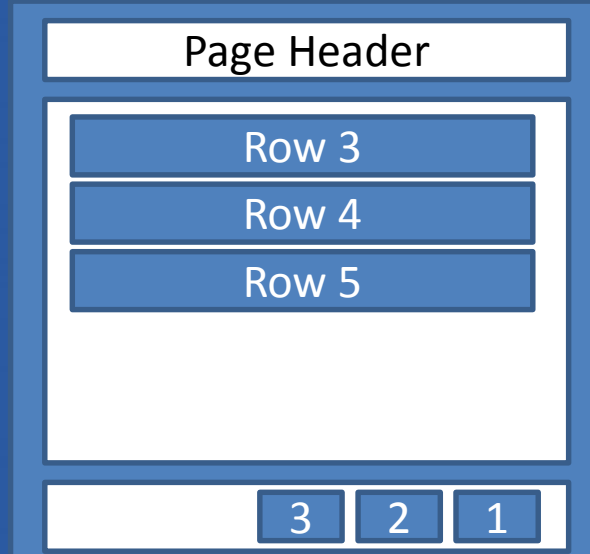
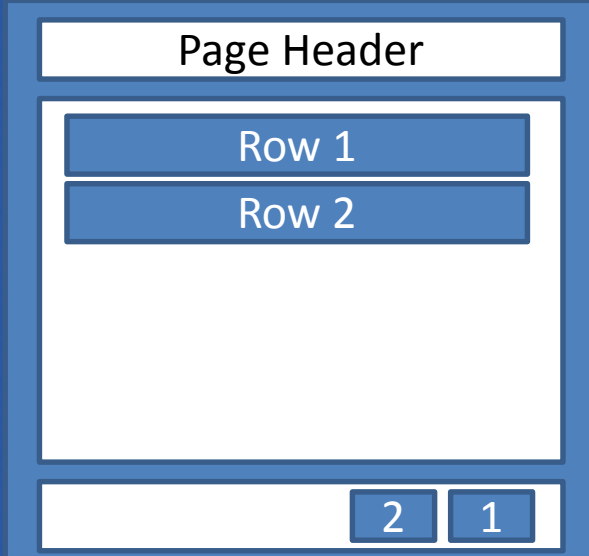
The problems:

- Size
 - Guids 16 bytes
 - Int 4 bytes
 - Disk and Memory
- Fragmentation

What Happens When a Page is Full



Fragmentation



GUID Primary Keys

A partial solution:

- NEWSEQUENTIALID()
 - Creates a sequential GUID.
 - Minimizes fragmentation
 - Still have the space issue.

Surrogate Key/No Alternate Key

- A surrogate key is a primary key automatically generated
- Surrogate keys are good thing.
- BUT when alternate key isn't defined bad things can happen...bad things man, bad things.
- Demo...

Suggestions

- Align proper design with the requirements.
- Learn how SQL Server works
- Work more closely with the DB team.
- Document DB design decisions.

Reference

- **Data Model Resource Book vol 1-3**
Len Silverston
- **Pro SQL Server 2012 Relational Database Design and Implementation**
Louis Davidson, Jessica Moss

Thank You!!

Richie Rump

@Jorriss

<http://jorriss.net>

<http://slideshare.net/jorriss>

<http://dotnetmiami.com>