

Introduction to Execution Plans

Kevan Riley



#372 | EXETER 2015

Who is this guy?

- Kevan Riley



kevan.riley@rileywaterhouse.co.uk



@kevriley

- Independent freelance SQL Server
- 15+ years SQL Server
- Moderator at ask.sqlservercentral.com



Agenda

- What are Execution Plans?
- Why are they important?
- How to view an Execution Plan
- How to read Graphical Execution Plans
- Execution Plan Operators
- Performance Tuning

Execution Plans

- Query optimizer
 - Cost based optimizer
 - Cardinality Estimator
 - Generate execution plans
 - Evaluate least cost based on Statistics
 - Estimated execution plan -> Plan cache
- Storage engine -> Actual execution plan

Execution Plans



- Optimizer finds 'best' plan = least cost plan in the shortest time
- Plan cache enables plan reuse
 - Performance gain
- Plan cache aged out / cleared

Statistics

- Data about your data
- Statistics on columns and indexes
 - Selectivity
 - Uniqueness
 - Distribution
- Automatically created and maintained by default
- Dave Morrison – “Statistics, Estimation & Plan Caching - The Big Three” next session in this room!

View Execution Plans

- SQL Server Management Studio (SSMS)



- SET statements
- DMVs (plan cache)
- Extended Events (SQL2008+ only)
- SQL Profiler

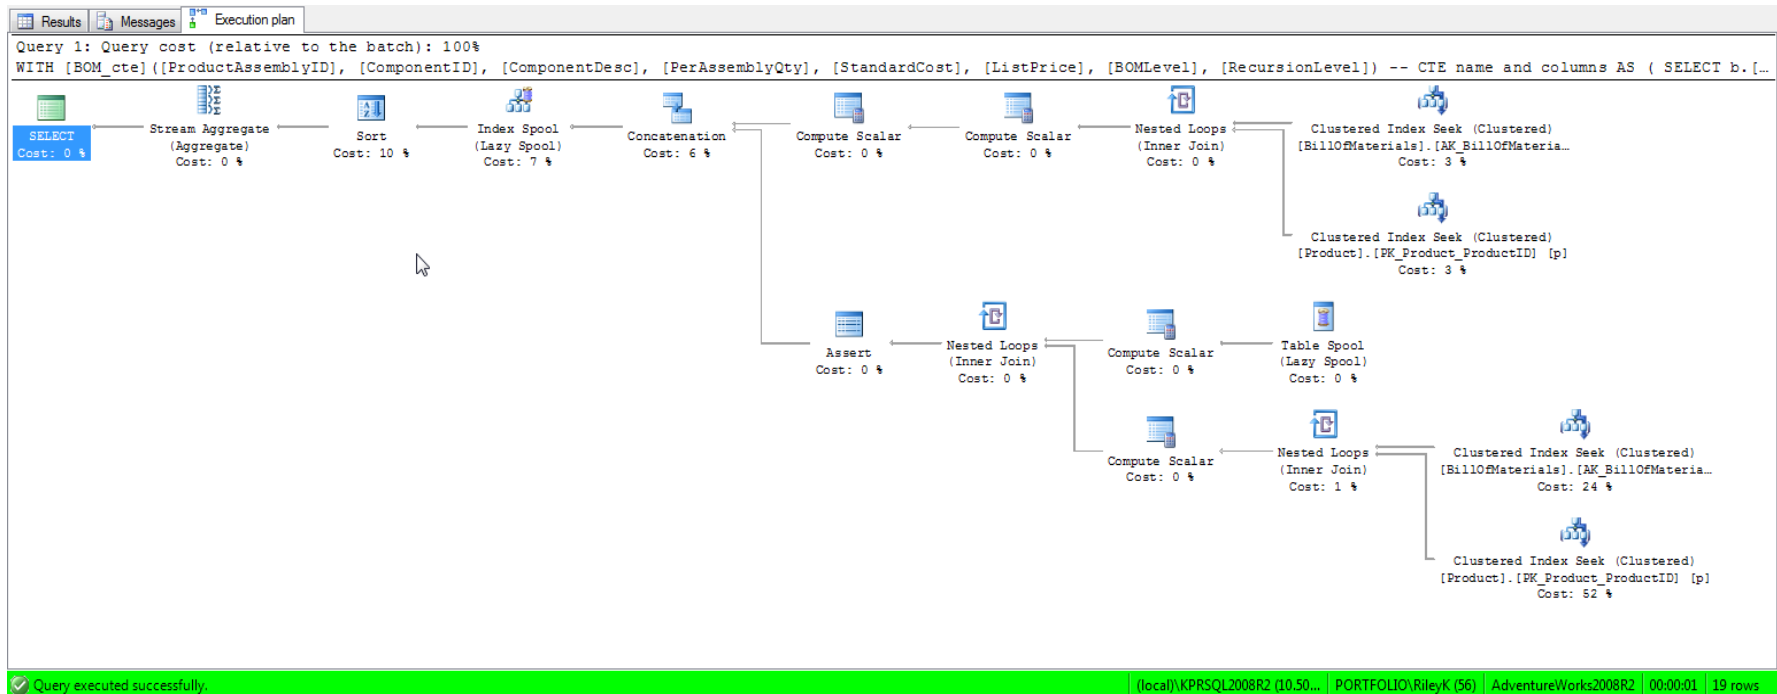
View Execution Plans - Graphical

- Estimated

- Ctrl-L
- Menu bar

- Actual

- Ctrl-M
- Menu Bar



View Execution Plans - Text

■ Estimated

- SET SHOWPLAN_ALL ON
- SET SHOWPLAN_TEXT ON

■ Actual

- SET STATISTICS PROFILE ON

Rows	Executes	StmtText
19	1	WITH [BOM_cte] ([ProductAssemblyID], [ComponentID], [ComponentDesc], [PerAssemblyQty], [StandardCost], [ListPr:
		AS (
		SELECT b.[ProductAssemblyID], b.[ComponentID], p.[Name], b.[PerAssem 1 1 0 NULL NULL
19	1	--Stream Aggregate(GROUP BY:([Recr1030], [Recr1024], [Recr1025], [Recr1026], [Recr1031], [Recr1028], [Recr:
19	1	--Sort(ORDER BY:([Recr1030] ASC, [Recr1024] ASC, [Recr1025] ASC, [Recr1026] ASC, [Recr1031] ASC, [Recr:
19	1	--Index Spool(WITH STACK)
19	1	--Concatenation
0	0	--Compute Scalar(DEFINE:([Expr1034]=(0)))
0	0	--Compute Scalar(DEFINE:([Expr1007]=(0)))
8	1	--Nested Loops(Inner Join, OUTER REFERENCES:([b].[ComponentID]))
8	1	--Nested Loops(Left Anti Semi Join)
8	1	--Clustered Index Seek(OBJECT:([AdventureWorks2008R2].[Production].
0	8	--Row Count Spool
0	1	--Clustered Index Scan(OBJECT:([AdventureWorks2008R2].[Sales].
8	8	--Clustered Index Seek(OBJECT:([AdventureWorks2008R2].[Production].[Proc
11	1	--Assert(WHERE:(CASE WHEN [Expr1036]>(25) THEN (0) ELSE NULL END))
11	1	--Nested Loops(Inner Join, OUTER REFERENCES:([Expr1036], [Recr1008], [Recr1009],
0	0	--Compute Scalar(DEFINE:([Expr1036]=[Expr1035]+(1)))
19	1	--Table Spool(WITH STACK)
0	0	--Compute Scalar(DEFINE:([Expr1023]=[Recr1015]+(1)))
11	19	--Nested Loops(Inner Join, OUTER REFERENCES:([b].[ComponentID]))
11	19	--Nested Loops(Left Anti Semi Join)
11	19	--Clustered Index Seek(OBJECT:([AdventureWorks2008R2].[Product
0	11	--Row Count Spool
0	1	--Clustered Index Scan(OBJECT:([AdventureWorks2008R2].[S:
11	11	--Clustered Index Seek(OBJECT:([AdventureWorks2008R2].[Production].

View Execution Plans - XML

- Estimated
 - SET SHOWPLAN_XML ON
- Actual
 - SET STATISTICS XML ON

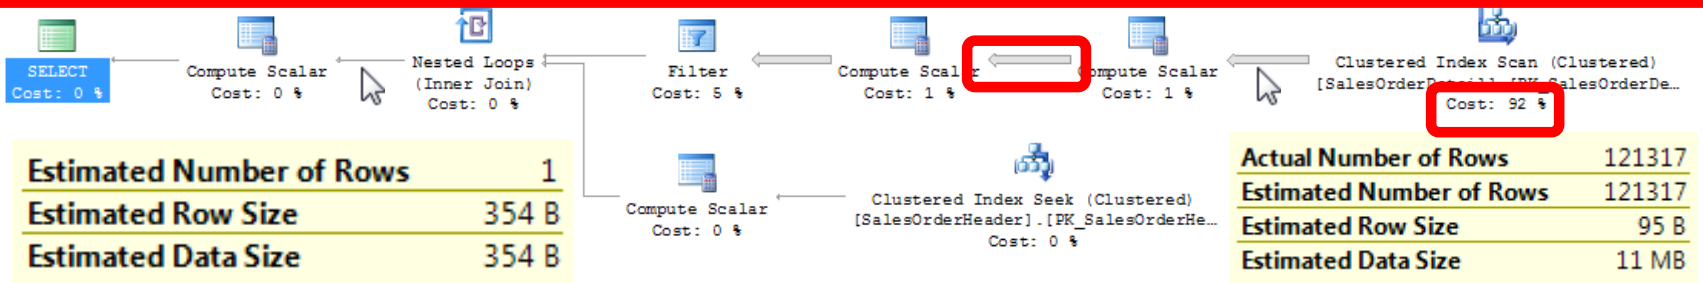
Execution plan.xml | Disk Usage by T...45\KPRSQL2008R2 | SQLQuery15.sql* | HEAPS.sql | SQLQuery13.sql* | SQLQuery10.sql* | SQLQuery9.sql*

```
<?xml version="1.0" encoding="utf-16"?>
<ShowPlanXML xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance" xmlns:xsd="http://www.w3.org/2001/XMLSchema" Version="1.1" Build="10.50.4000.0" xmlns="ht
<BatchSequence>
  <Batch>
    <Statements>
      <StmtSimple StatementCompId="4" StatementEstRows="9.18693" StatementId="1" StatementOptmLevel="FULL" StatementSubTreeCost="2.42492" StatementText="W
      <StatementSetOptions ANSI_NULLS="true" ANSI_PADDING="true" ANSI_WARNINGS="true" ARITHABORT="true" CONCAT_NULL_YIELDS_NULL="true" NUMERIC_ROUNDABOR
      <QueryPlan DegreeOfParallelism="1" MemoryGrant="1024" CachedPlanSize="80" CompileTime="447" CompileCPU="354" CompileMemory="1256">
        <RelOp AvgRowSize="108" EstimateCPU="1.01056E-05" EstimateIO="0" EstimateRebinds="0" EstimateRewinds="0" EstimateRows="9.18693" LogicalOp="Aggre
        <OutputList>
          <ColumnReference Column="Recr1024" />
          <ColumnReference Column="Recr1025" />
          <ColumnReference Column="Recr1026" />
          <ColumnReference Column="Recr1028" />
          <ColumnReference Column="Recr1029" />
          <ColumnReference Column="Recr1030" />
          <ColumnReference Column="Recr1031" />
          <ColumnReference Column="Expr1032" />
        </OutputList>
        <RunTimeInformation>
          <RunTimeCountersPerThread Thread="0" ActualRows="19" ActualEndOfScans="1" ActualExecutions="1" />
        </RunTimeInformation>
        <StreamAggregate>
          <DefinedValues>
            <DefinedValue>
              <ColumnReference Column="Expr1032" />
              <ScalarOperator ScalarString="SUM([Recr1027])">
                <Aggregate AggType="SUM" Distinct="false">
```

How to read Execution Plans

■ Right to left ?

Query 1: Query cost (relative to the batch) : 100%
select * from Sales.SalesOrderDetail join Sales.SalesOrderHeader on Sales.SalesOrderHeader.SalesOrderID = Sales.SalesOrderDetail.Sale
Missing Index (Impact 56.1328): CREATE NONCLUSTERED INDEX [<Name of Missing Index, sysname,>] ON [Sales].[SalesOrderDetail] ([OrderQt



- Batch cost
- Operator costs
- Arrows
- Missing Indexes

How to read Execution Plans

■ Properties

Index Scan (NonClustered)	
Scan a nonclustered index, entirely or only a range.	
Physical Operation	Index Scan
Logical Operation	Index Scan
Actual Number of Rows	19972
Estimated I/O Cost	0.0794213
Estimated CPU Cost	0.0221262
Estimated Number of Executions	1
Number of Executions	1
Estimated Operator Cost	0.101547 (22%)
Estimated Subtree Cost	0.101547
Estimated Number of Rows	19972
Estimated Row Size	117 B
Actual Rebinds	0
Actual Rewinds	0
Ordered	False
Node ID	2
Object	
[AdventureWorks2008R2].[Person].[Person]. [IX_Person_LastName_FirstName_MiddleName]	
Output List	
[AdventureWorks2008R2].[Person]. [Person].BusinessEntityID, [AdventureWorks2008R2]. [Person].[Person].FirstName, [AdventureWorks2008R2].[Person].[Person].LastName	

Properties	
Index Scan (NonClustered)	
Misc	
Actual Number of Rows	19972
Actual Rebinds	0
Actual Rewinds	0
Defined Values	[AdventureWorks2008R2].[Pers
Description	Scan a nonclustered index, enti
Estimated CPU Cost	0.0221262
Estimated I/O Cost	0.0794213
Estimated Number of Exec	1
Estimated Number of Row	19972
Estimated Operator Cost	0.101547 (22%)
Estimated Rebinds	0
Estimated Rewinds	0
Estimated Row Size	117 B
Estimated Subtree Cost	0.101547
Forced Index	False
ForceScan	False
ForceSeek	False
Logical Operation	Index Scan
Node ID	2
NoExpandHint	False
Number of Executions	1
Object	[AdventureWorks2008R2].[Pers
Ordered	False
Output List	[AdventureWorks2008R2].[Pers
Parallel	False
Physical Operation	Index Scan
TableCardinality	19972

Operators

- Data retrieval
 - Table Scan
 - Clustered Index Scan
 - Clustered Index Seek
 - Non-clustered Index Scan
 - Non-clustered Index Seek
 - Key Lookup
 - RID Lookup
- Join (loop/merge/hash)
- Sort

Examples given tested against:

- ✓ *AdventureWorks2008R2*
- ✓ *AdventureWorks2012*
- ✓ *AdventureWorks2014*

Data Operations – Table Scans

```
select *  
from Production.ProductProductPhoto
```



Table Scan
[ProductProductPhoto]
Cost: 100 %

- Table Scan – Heap
 - All (or majority) of rows
 - No useful indexes
 - Small tables
 - Table variables / CTE
- Issue?
 - Large number of rows

Data Operations – Index Scans

```
select *  
from Sales.Currency
```



```
Clustered Index Scan (Clustered)  
[Currency].[PK_Currency_CurrencyCod..  
Cost: 100 %
```

- Clustered Index Scans
 - Same as a table scan !
- Issue?
 - Large range of data is been selected
 - Maybe benefit from different index
 - Stale statistics

Data Operations – Index Seeks

```
select CustomerID, PersonId  
from Sales.Customer  
where CustomerID < 10
```



```
Clustered Index Seek (Clustered)  
[Customer].[PK_Customer_CustomerID]  
Cost: 100 %
```

- Clustered Index Seek
 - ‘Gold’ standard
 - All data available at leaf level

Data Operations – Index Scans

```
select BusinessEntityID,  
       FirstName, LastName  
from Person.Person
```



```
Index Scan (NonClustered)  
[Person].[IX_Person_LastName_FirstN..  
Cost: 100 %
```

- Non-Clustered Index Scan
 - No optimal index for the query
 - Dataset returned represents most of the table
- Issue?
 - Refine WHERE clause
 - Lookups (coming later....)

Data Operations – Index Seeks

```
select BusinessEntityID, FirstName, LastName  
from Person.Person  
where LastName = 'Riley'
```




```
Index Seek (NonClustered)  
[Person].[IX_Person_LastName_FirstN...  
Cost: 100 %
```

- Non-Clustered Index Seek
 - Similar to CIX seek but only index fields available
- Issue?
 - May cause Lookups

Data Operations - Lookups

- Only on non-clustered index operations
- RID Lookup (Heap)

```
select DatabaseLogID, DatabaseUser  
from dbo.DatabaseLog  
where DatabaseLogID = 100
```



Data Operations - Lookups

- Only on non-clustered index operations
- Key Lookup

```
select BusinessEntityID, FirstName, LastName, PersonType  
from Person.Person  
where LastName = 'Riley'
```



Index Seek (NonClustered)
[Person].[IX_Person_LastName_FirstN...
Cost: 47 %



Key Lookup (Clustered)
[Person].[PK_Person_BusinessEntityI...
Cost: 53 %

Data Operations - Lookups



— RID Lookup (Heap)
[DatabaseLog]
Cost: 50 %



Key Lookup (Clustered)
[Person].[PK_Person_BusinessEntityL...]
Cost: 53 %

- Eliminating lookups
 - Review column selection
 - Covering indexes (key fields or INCLUDEs)
- Exceptions
 - Select every column from the table

Join Operations

- Three main logical join operators
 - Nested Loop join
 - Merge join
 - Hash join

Join Operation – Nested Loop

- Very efficient operation
- Takes an outer data set and compares one row at a time to inner data set
- Most effective when
 - Outer input is small
 - Inner (small or large) is indexed



Nested Loops
(Inner Join)
Cost: 0 %

Join Operations - Merge



- Sorted data makes this a very efficient operation
- Addition of Sort operator makes it less efficient
- Otherwise Hash Join

┆ Merge Join
(Inner Join)
Cost: 25 %

Join Operations – Hash Join



- Large, unsorted data
- Efficient where no useable indexes

Hash Match
(Inner Join)
Cost: 67 %

Join Operations - Summary

- No 'ideal' join – depends on data
- Join operators even if there is no JOIN statement
- Resource-wise Nested Loop is 'best'
- Merge requires sorted data – can be forced

Sort

- Orders the data
 - ORDER BY
 - Ranking functions
 - ROW_NUMBER()
 - RANK()
 - DENSE_RANK()
 - NTILE()
 - Merge Join
- Blocking Operator
- Can be expensive



Sort
Cost: 22 %

Performance considerations


- Seek vs. Scan
 - Missing indexes
 - Statistics
 - Functions on indexed columns non-Sargable
- Actual rows vs. Estimated Rows
 - Statistics
 - Table variables (consider temp tables)
 - User defined functions
 - Scalar functions
 - Multi-statement table valued functions
 - Consider re-writing to inline table valued functions

Performance considerations

- Sort operations expensive
 - Statistics are wrong – spill to tempdb/disk
- Correct join operators
 - Statistics
- Review Lookups
- Did I mention statistics?

Further Resources

- SQLSentry Plan Explorer FREE!
- SQL Server Execution Plans – Grant Fritchey
[@GFritchey](#) | www.scarydba.com
- Inside the SQL Server Query Optimizer – Benjamin Nevarez
[@BenjaminNevarez](#) | www.benjaminnevarez.com



Takeaway

- Execution plan is your window
- Get familiar
- Statistics!

Thanks for listening



kevan.riley@rileywaterhouse.co.uk
www.rileywaterhouse.co.uk
[@kevriley](https://twitter.com/kevriley)

Today is brought to you by

redgate
ingeniously simple


and in association with



Please visit our sponsors

SanDisk®

**BIG BANG
DATA CO**



Software



PYRAMID
ANALYTICS



Thanks for listening



kevan.riley@rileywaterhouse.co.uk
www.rileywaterhouse.co.uk
[@kevriley](https://twitter.com/kevriley)