Inside the Query Optimizer

Speaker Introduction
Bradley Ball

- Almost 15 Years IT Experience
- Previous experience DBA, for the U.S. Army, The Executive Office of the President, Sr. SQL DBA Staff Specialist at Publix
- Currently the Data Platform Management Lead for Pragmatic Works
- Microsoft VTSP for the Greater North East
- MCITP SQL 2005 DBA & SQL 2008 DBA
- Blog: http://www.SQLBalls.com
- Twitter: @SQLBalls @BradleyBall_PW
- Email: bball@PragmaticWorks.com
- Pro SQL Server 2012 Practices Author
- Chapter 14 PAGE & ROW COMPRESSION!

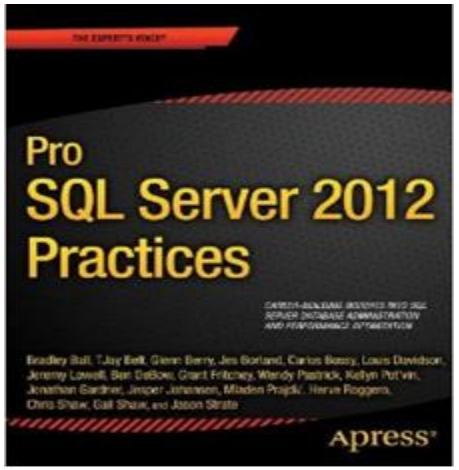

[Session Code]

Agenda

- Input Trees
- Optimize
- Joins
- Reasons to Join
- Hints, Plan Guides

FUN!

No Seriously FUN!

It won't be Lame I Promise

Why Optimize?

- You Want to go FAST!
- Everybody's Doing it!
- My Boss Told Me To

How T-SQL is Processed

How We Write SQL
SELECT TOP DISTINCT FROM JOIN ON WHERE GROUP BY CUBE / ROLLUP ORDER BY

SQL Reads What We Write
FROM ON JOIN WHERE GROUP BY CUBE /ROLLUP HAVING SELECT DISTINCT ORDER BY TOP

Relational Engine Query Process

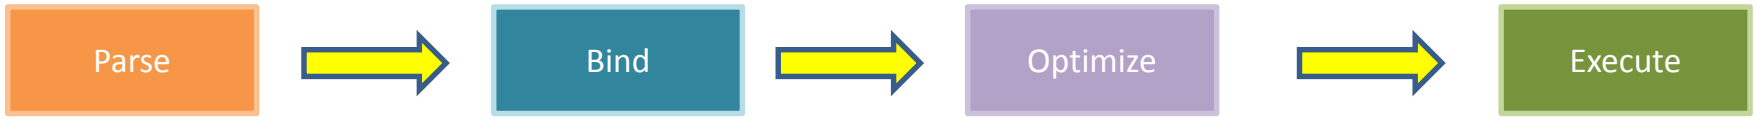

Relational Engine Query Process

Parse	Bind	Optimize	Execute

Parse
Validation of the syntax
No Executing a table

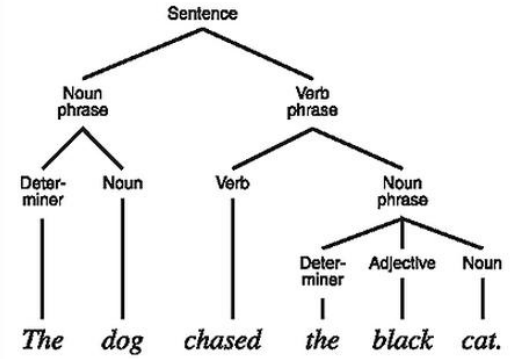

Relational Engine Query Process

Parse | Bind | Optimize | Execute

Bind

Name Resolution

Type Derivation

Aggregate Binding

Group Binding

Relational Engine Query Process

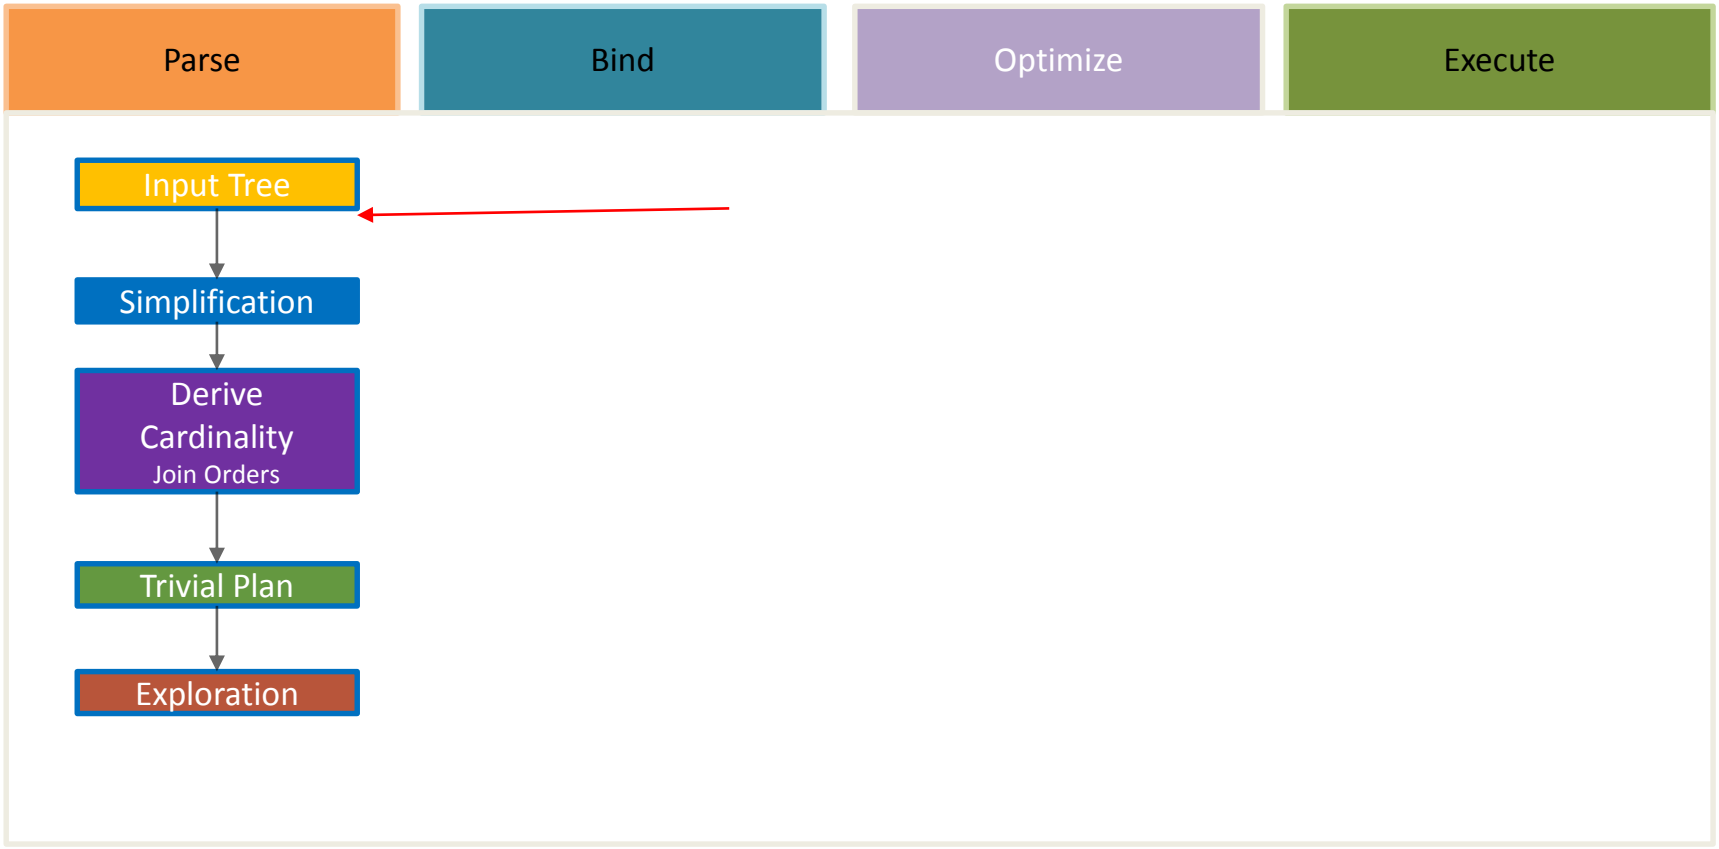

Relational Engine Query Process

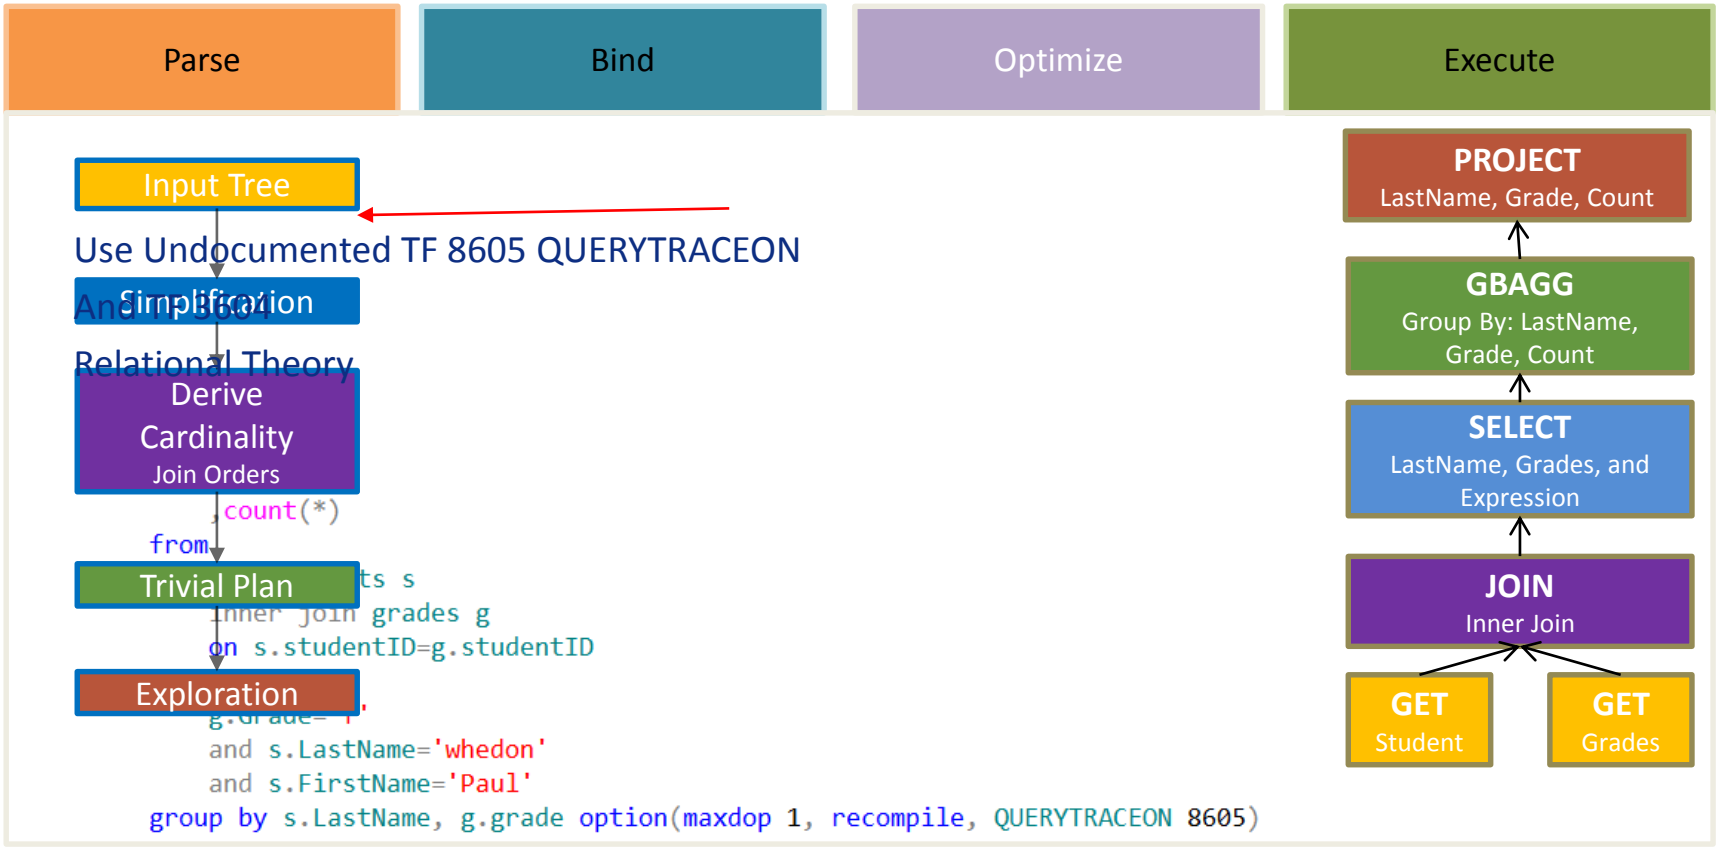

Relational Engine Query Process

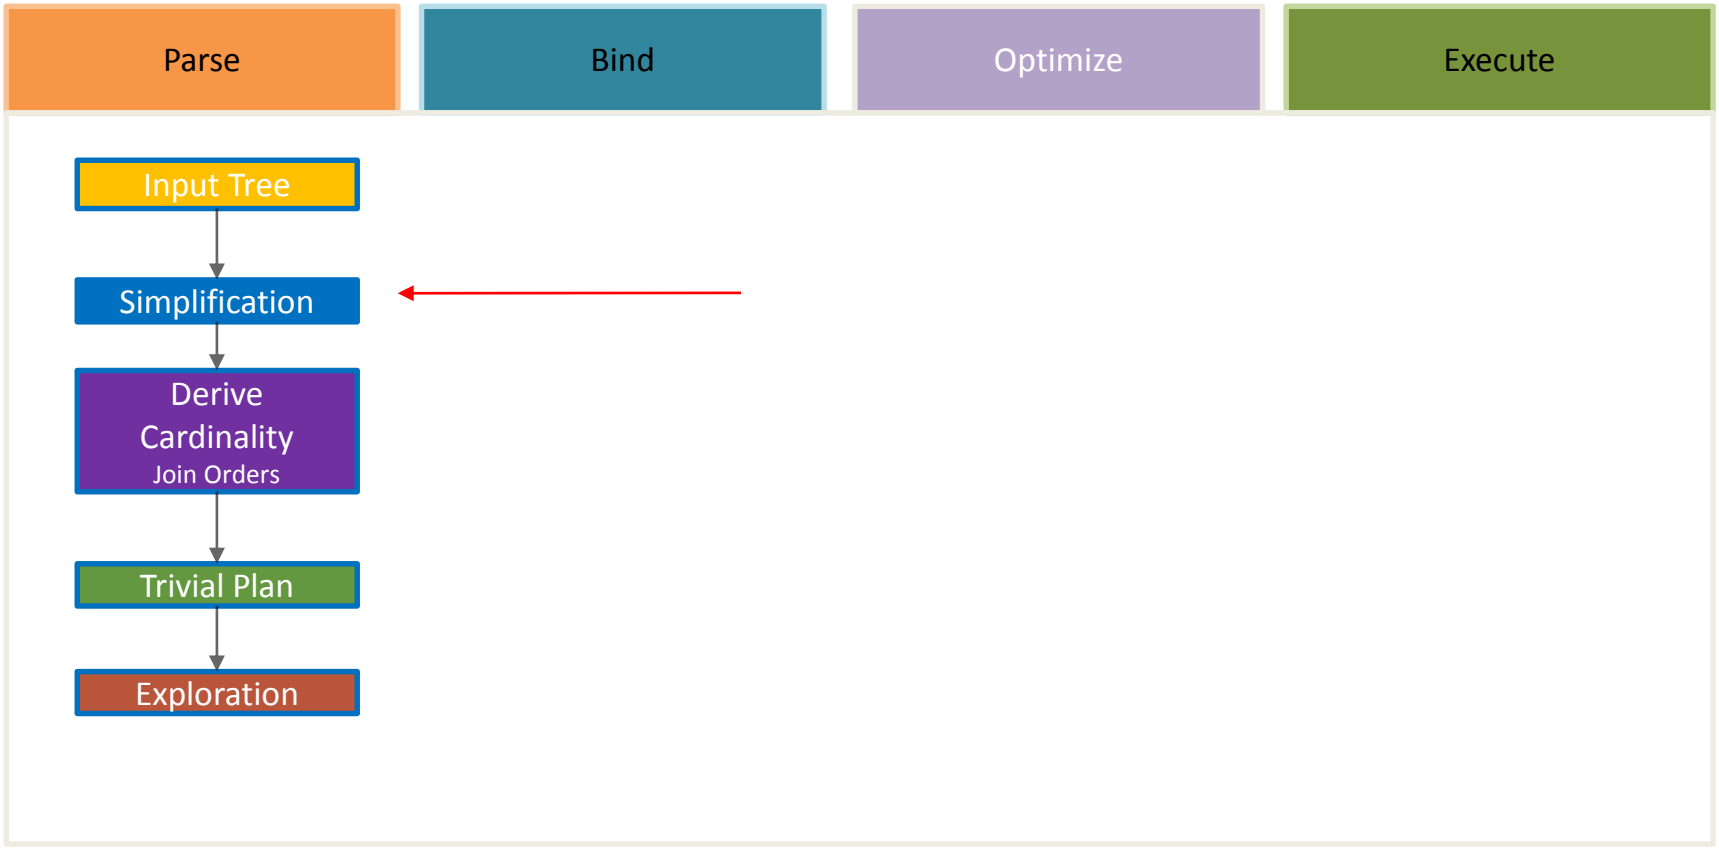

Relational Engine Query Process

Parse | Bind | Optimize | Execute

Simplification

- Constant Folding
  - Evaluating an Expression during optimization instead of having to constantly do it during Execution
- Domain Simplification
  - Allows the Optimizer to "Reason" the Valid Values within a Range
- Predicate push-down
  - Read Rob Farley, eliminate rows up front
- Join Simplification
  - Removes Unnecessary Joins (Powerful option)
- Contradiction Detection
  - Prevents I/O Reads when predicates contradict one another

Relational Engine Query Process

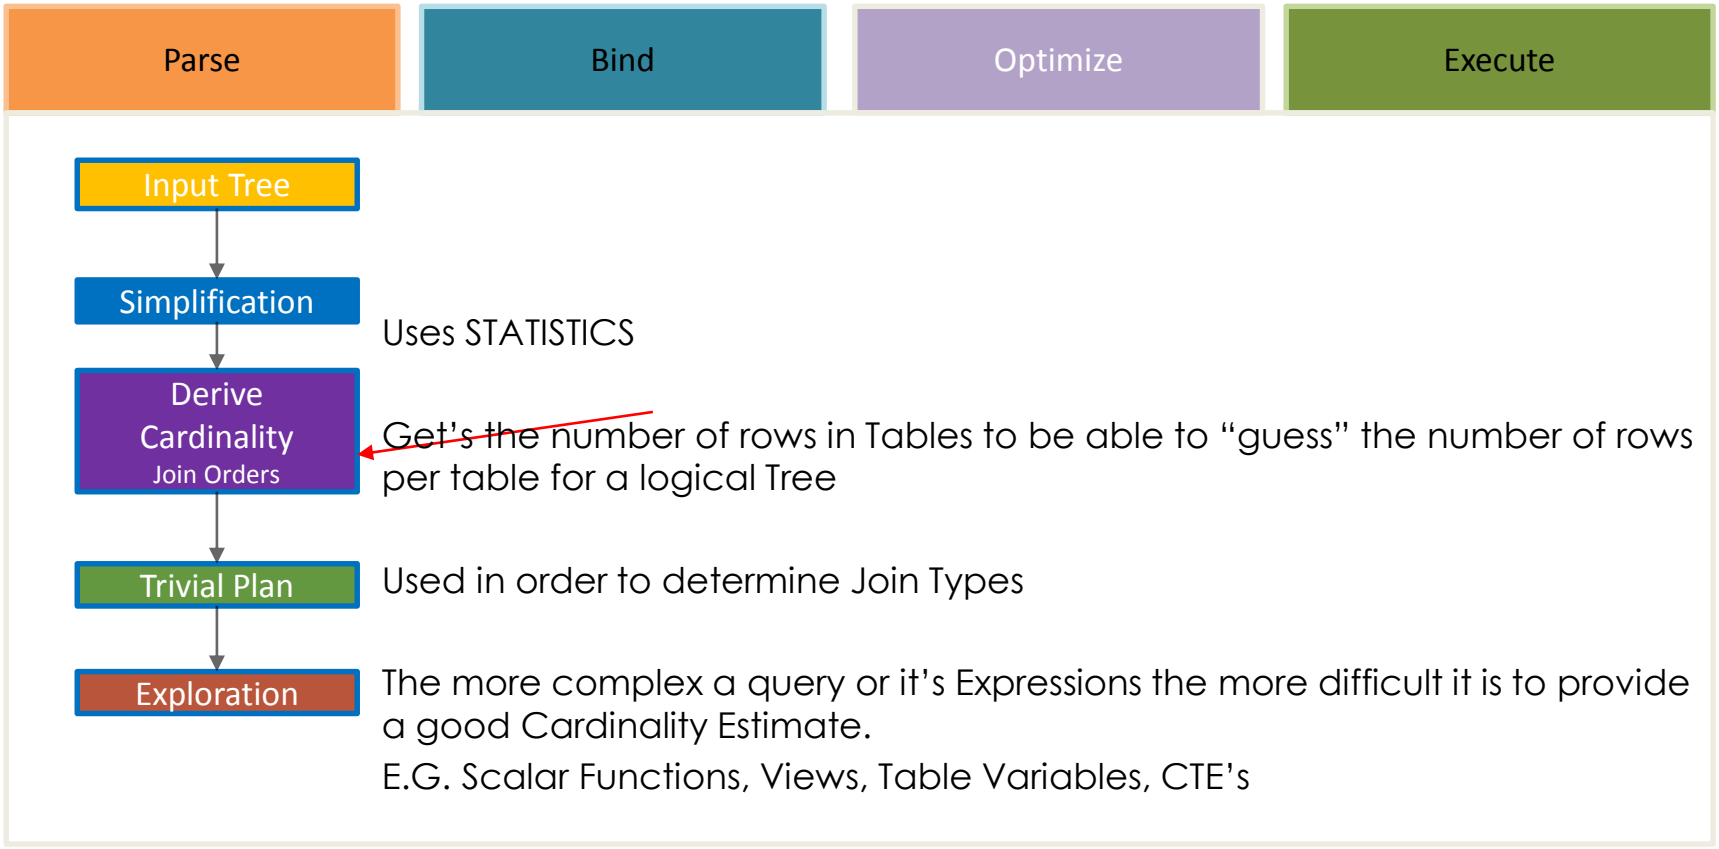

Relational Engine Query Process

Parse	Bind	Optimize	Execute

Derive Cardinality
Join Orders

Statistics

- Auto Update Statistics
- 20% of table plus 500 rows
- Not bad for small tables, but Really Bad for VLDB's
- Filtered Indexes Require the Same Base Table Updates as Other Non-Clustered Indexes

Relational Engine Query Process

Parse	Bind	Optimize	Execute

Derive Cardinality
Join Orders

Statistics on VLDB's

Trace Flag 2371

Table must have more than 25,000 Rows

Lowers the Threshold for Updates

100,000 = 10%
1,000,000 = 3.2%
10,000,000=1%
50,000,000=0.5%

Relational Engine Query Process

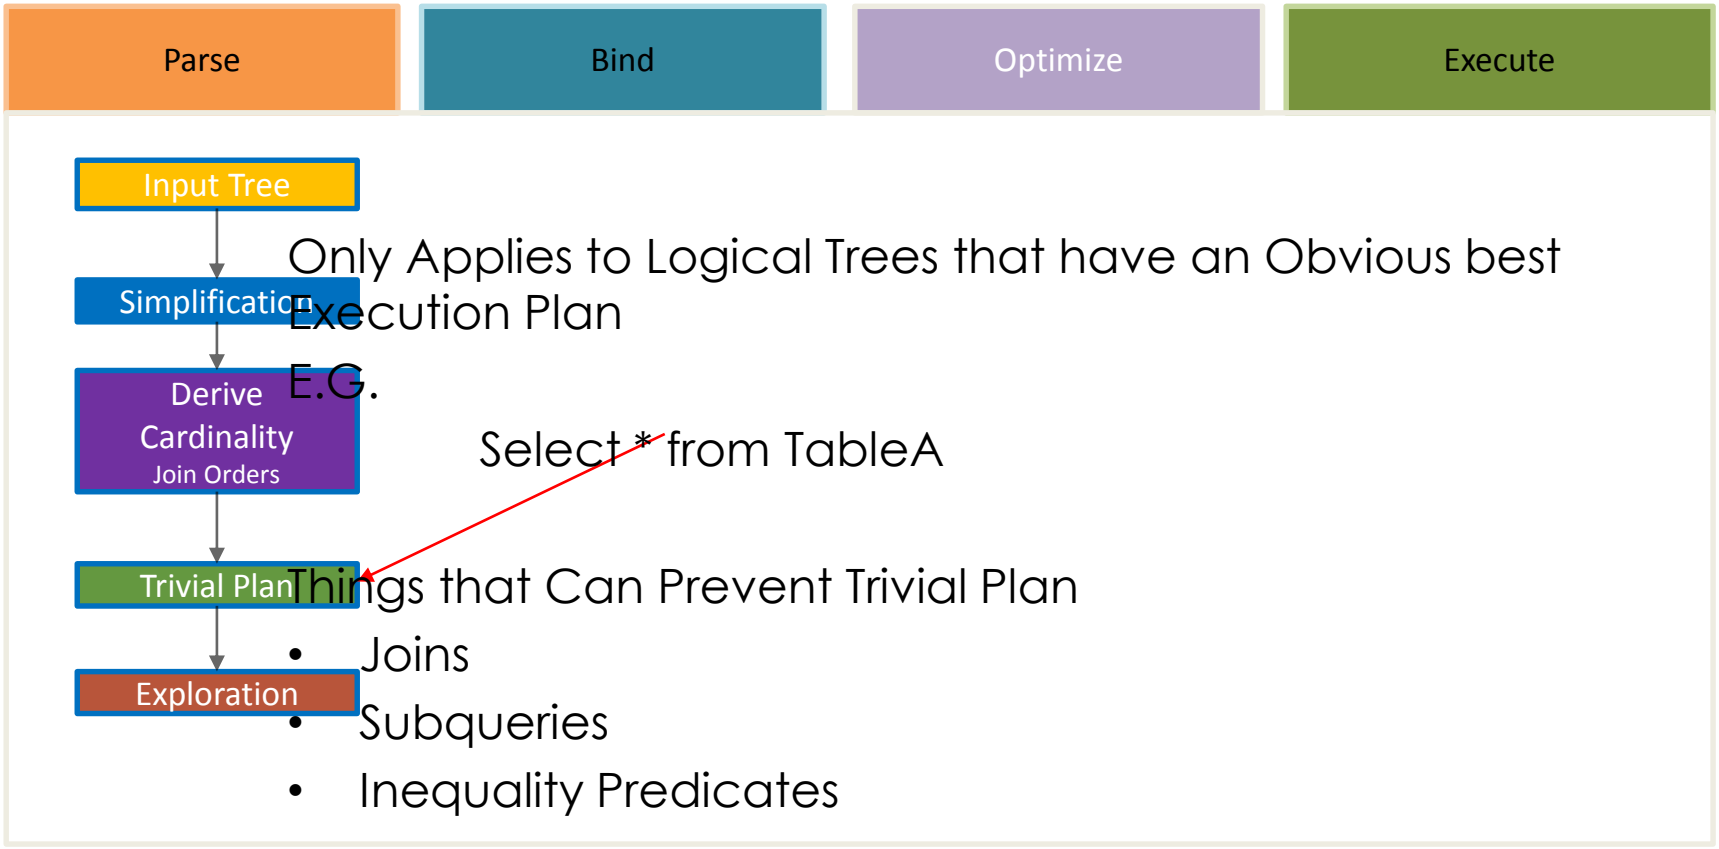

Relational Engine Query Process

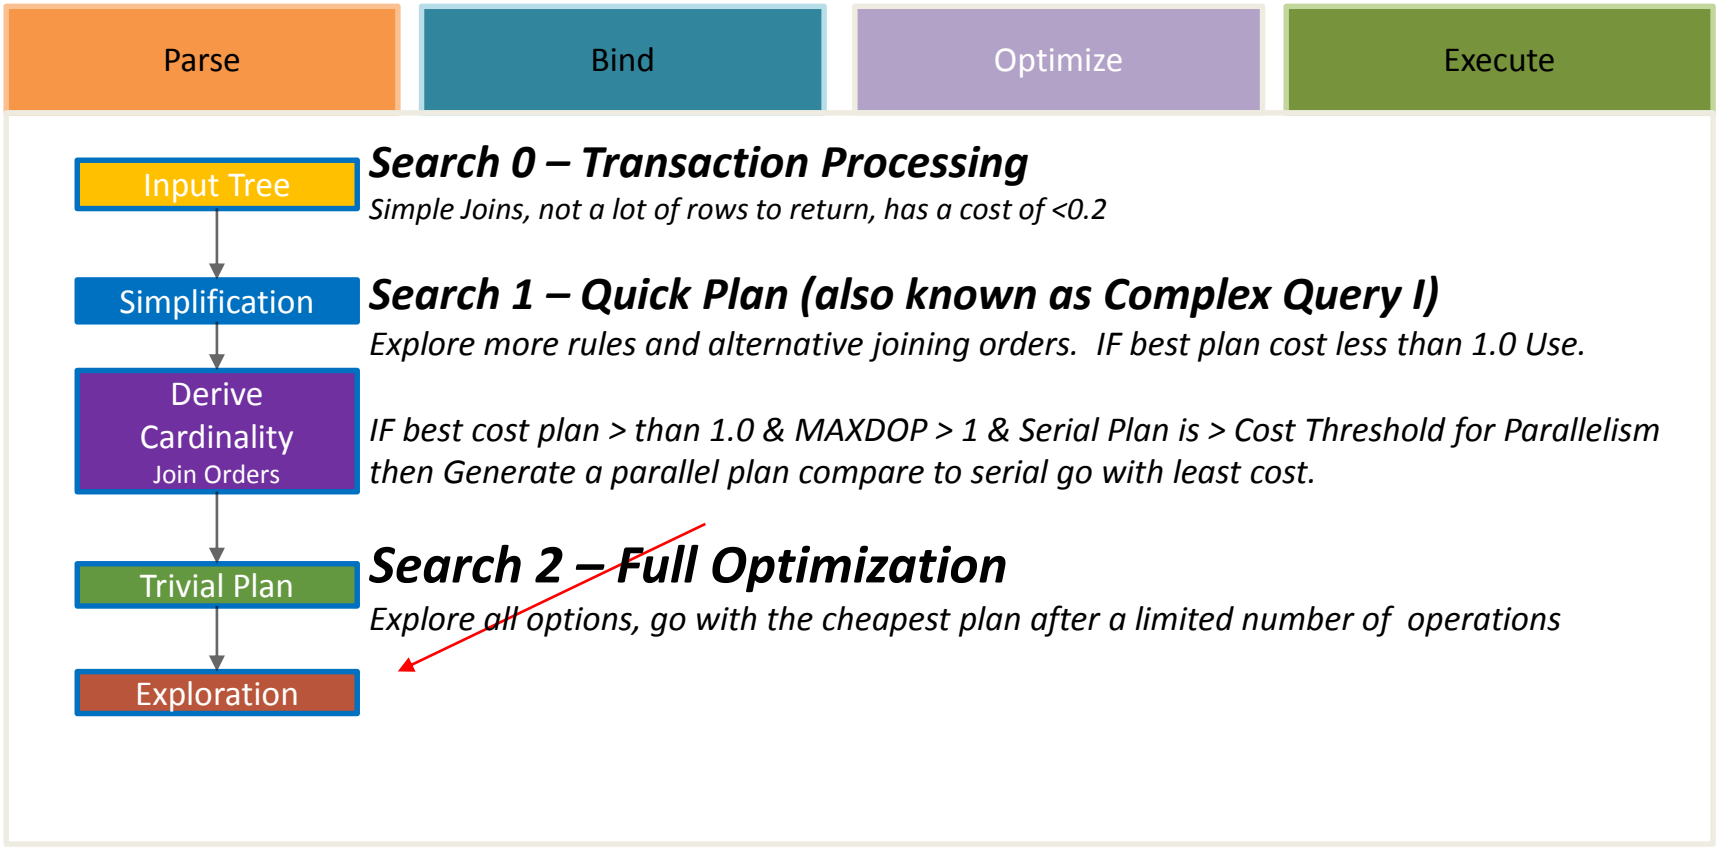

Relational Engine Query Process

Parse	Bind	Optimize	Execute

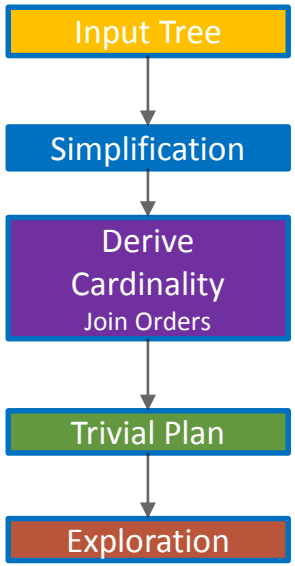

- Parse
- Algebrize
- Expand Views
- NNF convert
- Simplify
- Auto stats create /update
- Cardinality derivation
- Heuristic join reordering
- Project normalization
- Trivial Plan
- Exploration
- Copy Out
- Convert to executable Tree

Relational Engine Query Process

Parse	Bind	Optimize	Execute

Execution

- Physical Plan is Created
- The Memo is populated
- Plan is placed in the Plan Cache
- Interaction with the Storage Engine

Demo

What is Parallelism

A Process split across multiple CPU's/Cores

CAN speed up a plan

Optimizer determines Parallelism use based on

IF(Cost of Parallel Plan/Max DOP + Cost of Exchange Operators) < Cost of Serial Plan

What is Parallelism

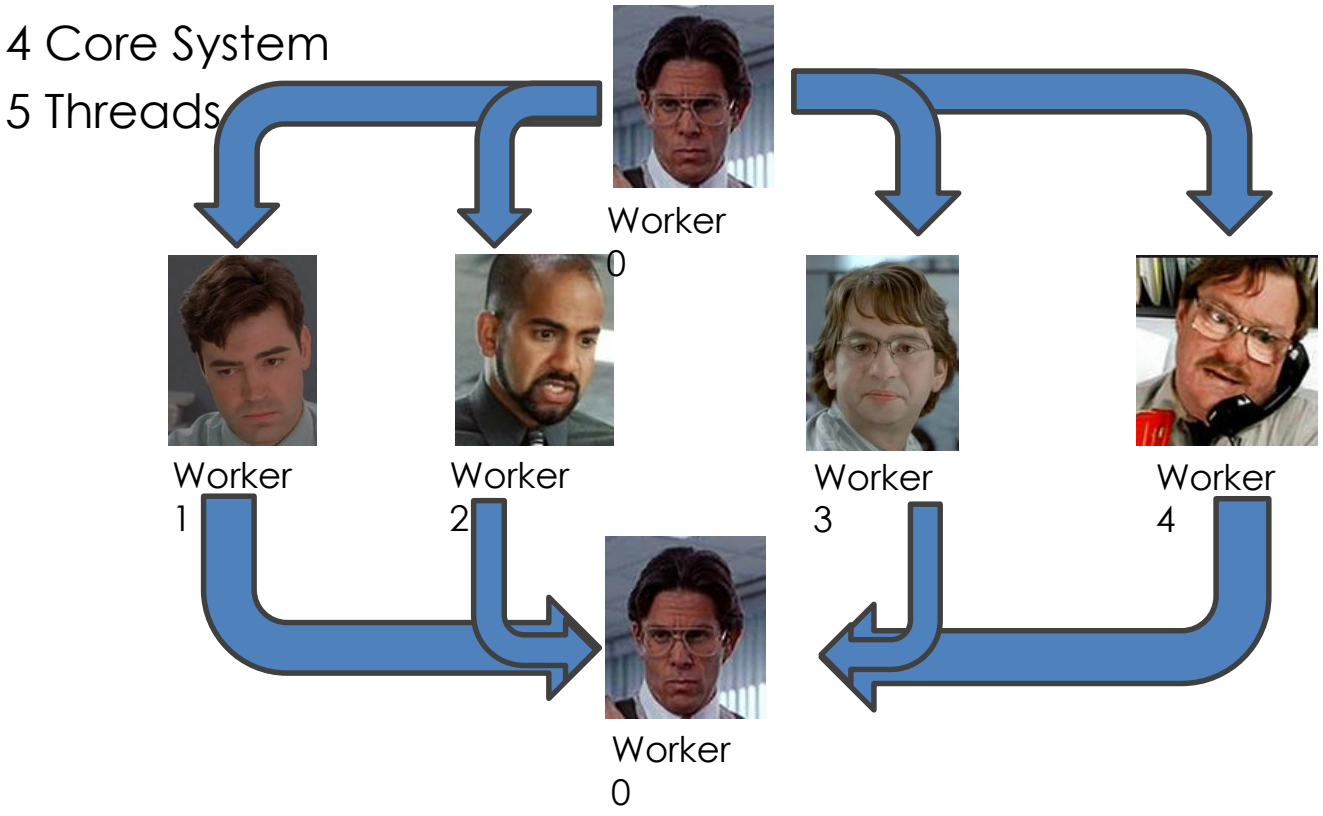

Cost Threshold for Parallelism

Default is 5

Has been 5 since 1997

Was the benchmark for someone on the Optimizer team at Microsoft.

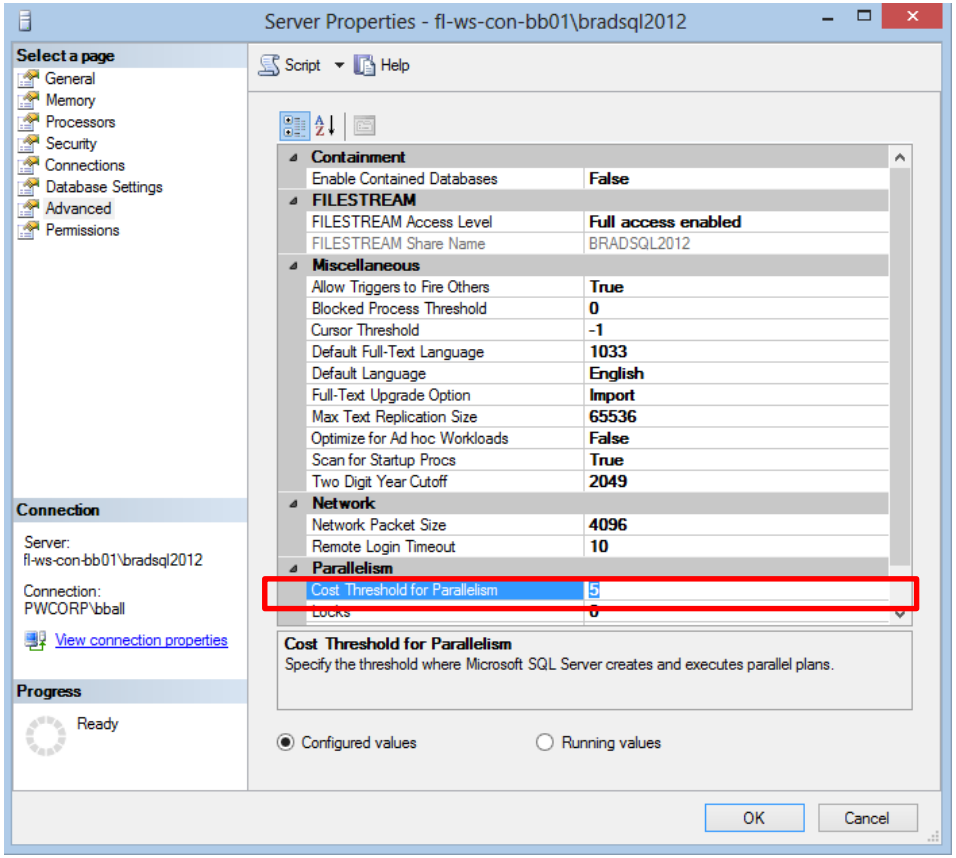

Demo

Syntactical Joins

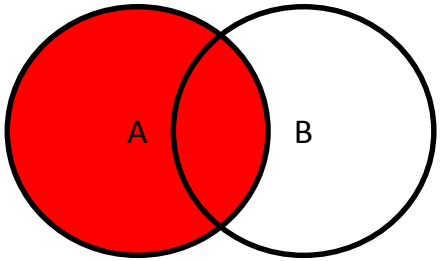


```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
On A.Key=B.Key
```

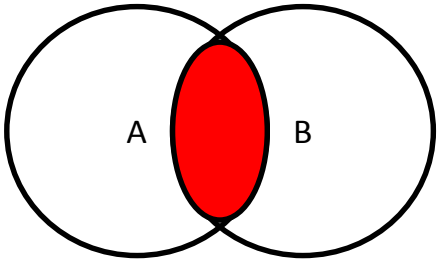


```
SELECT <select_list>
FROM TableA A
INNER JOIN TableB B
On A.Key=B.Key
```

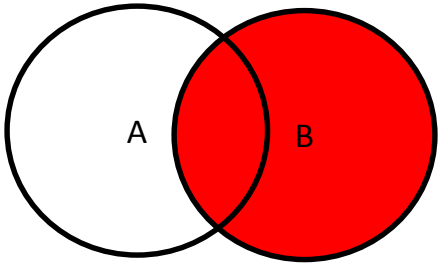


```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
On A.Key=B.Key
```

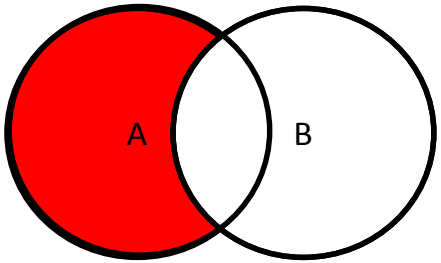


```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
On A.Key=B.Key
WHERE B.Key IS NULL
```

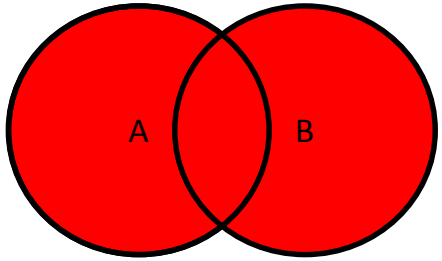


```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
On A.Key=B.Key
```

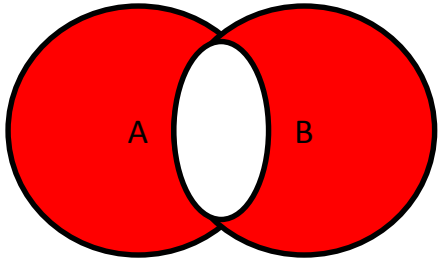


```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
On A.Key=B.Key
Where A.Key IS NULL OR B.Key IS
NULL
```

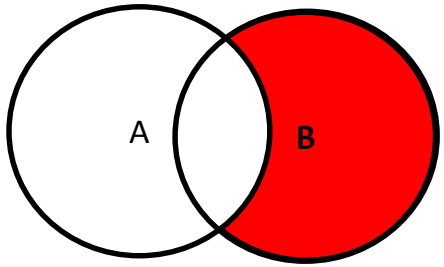


```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
On A.Key=B.Key
WHERE A.KEY IS NULL
```

Physical Operators

Joins

- Syntactical Joins
  - Inner Join
  - Outer Join
  - Cross Join
  - Cross Apply
  - Outer Apply
  - Semi-Join
  - Anti-Semi Join
- Physical Joins
  - Nested Loop Join
  - Merge Join
  - Hash Join

***If no Syntactical Join Type is specified Inner Join is the default**

Nested Loop Join

Compares Each Row from the Outer Table To Each Row in the Inner Table

```
for each row R1 in the outer table
   begin
      for each row R2 in the inner table
         if R1 joins with R2
            return (R1, R2)
          if R1 did not join
             return (R1, NULL)

   end
```

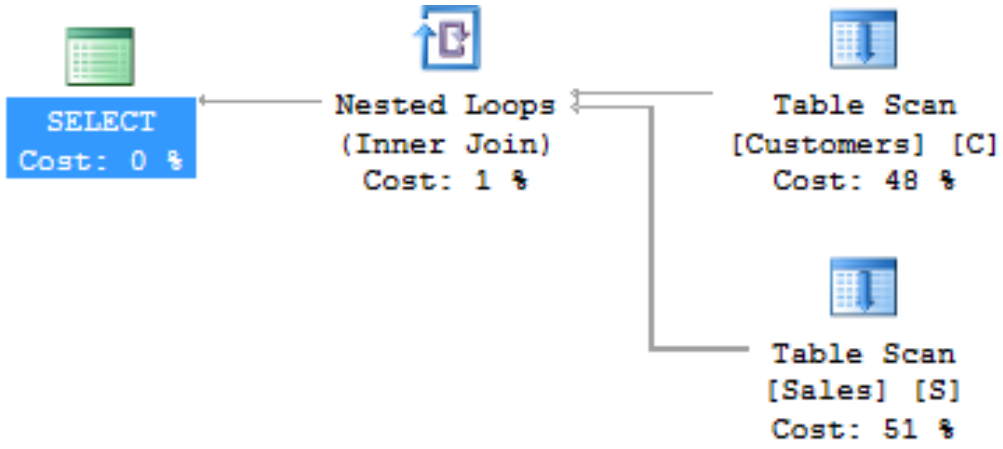

Sorted Merge Join

Simultaneously Reads and Compares two Sorted Inputs One Row at a Time

```
get first row R1 from input 1
get first row R2 from input 2
while not at the end of either input
   begin
      if R1 joins with R2
         begin
            return (R1, R2)
            get next row R2 from input 2
         end
      else if R1 < R2
         get next row R1 from input 1
      else
         get next row R2 from input 2
   end
```

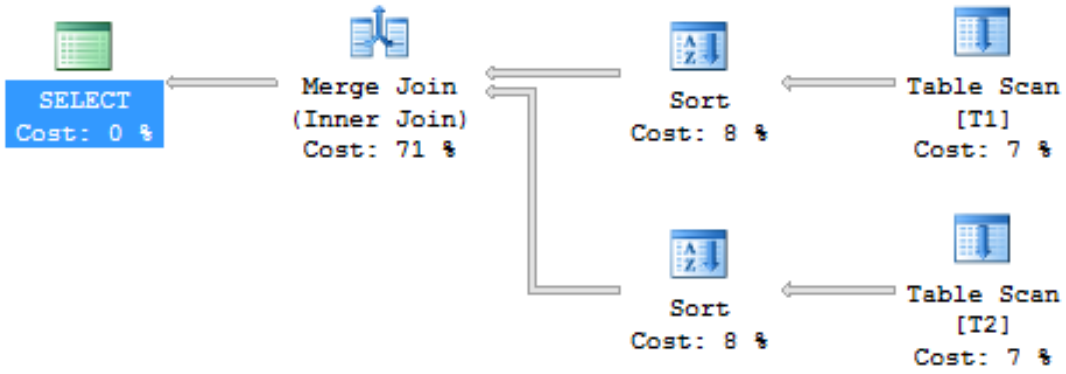

Hash Join

Join Heavy Lifter. Built in Two Phases: Build & Probe.

- Build
  - Reads Rows from 1st Input
  - Hashes on Equijoin Keys
  - Creates In-Memory Hash Table

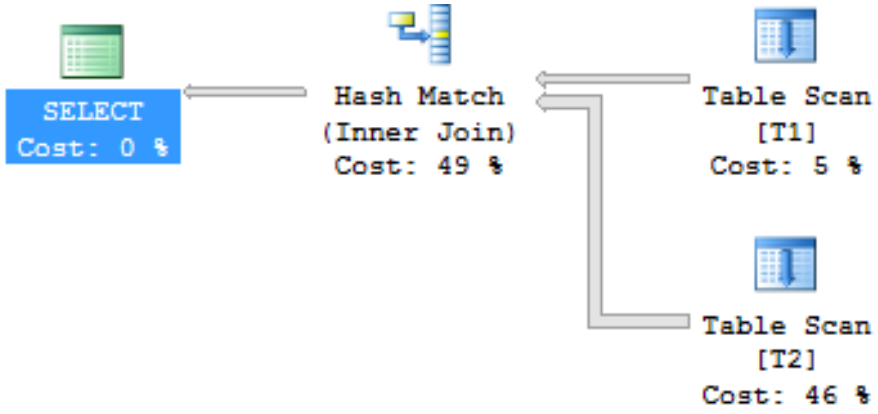

- Probe
  - Reads All Rows from 2nd Input
  - Hashes on Same Equijoin Keys
  - Probes for Matching Rows in Hash Table

```
for each row R1 in the build table
   begin
      calculate hash value on R1 join key(s)
      insert R1 into the appropriate hash bucket
   end
for each row R2 in the probe table
   begin
      calculate hash value on R2 join key(s)
      for each row R1 in the corresponding hash bucket
         if R1 joins with R2
            return (R1, R2)
   end
```

Hash Join

Has 3 Types of Tree Structures

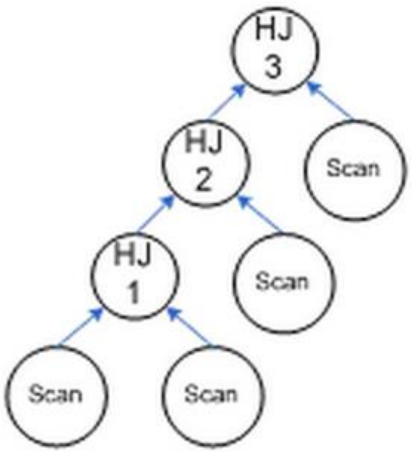

Left Deep

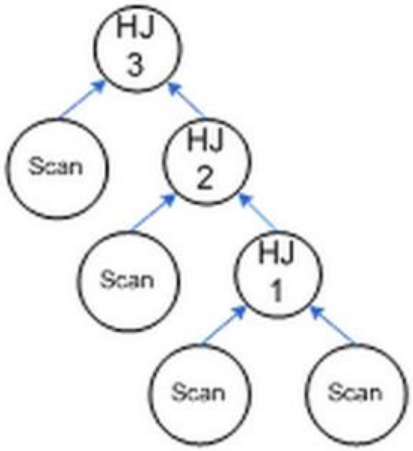

Right Deep

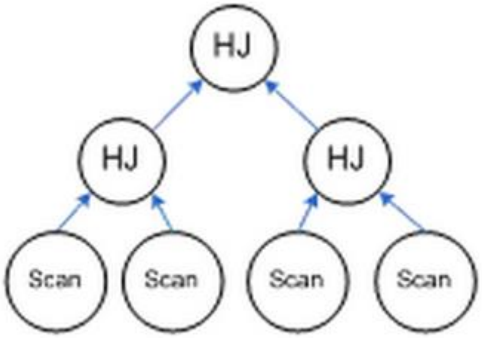

Bushy

Reasons to Join

- Extra Columns
- Eliminating Rows
- Duplicating Rows
- Introducing Nulls

Demo

Resources

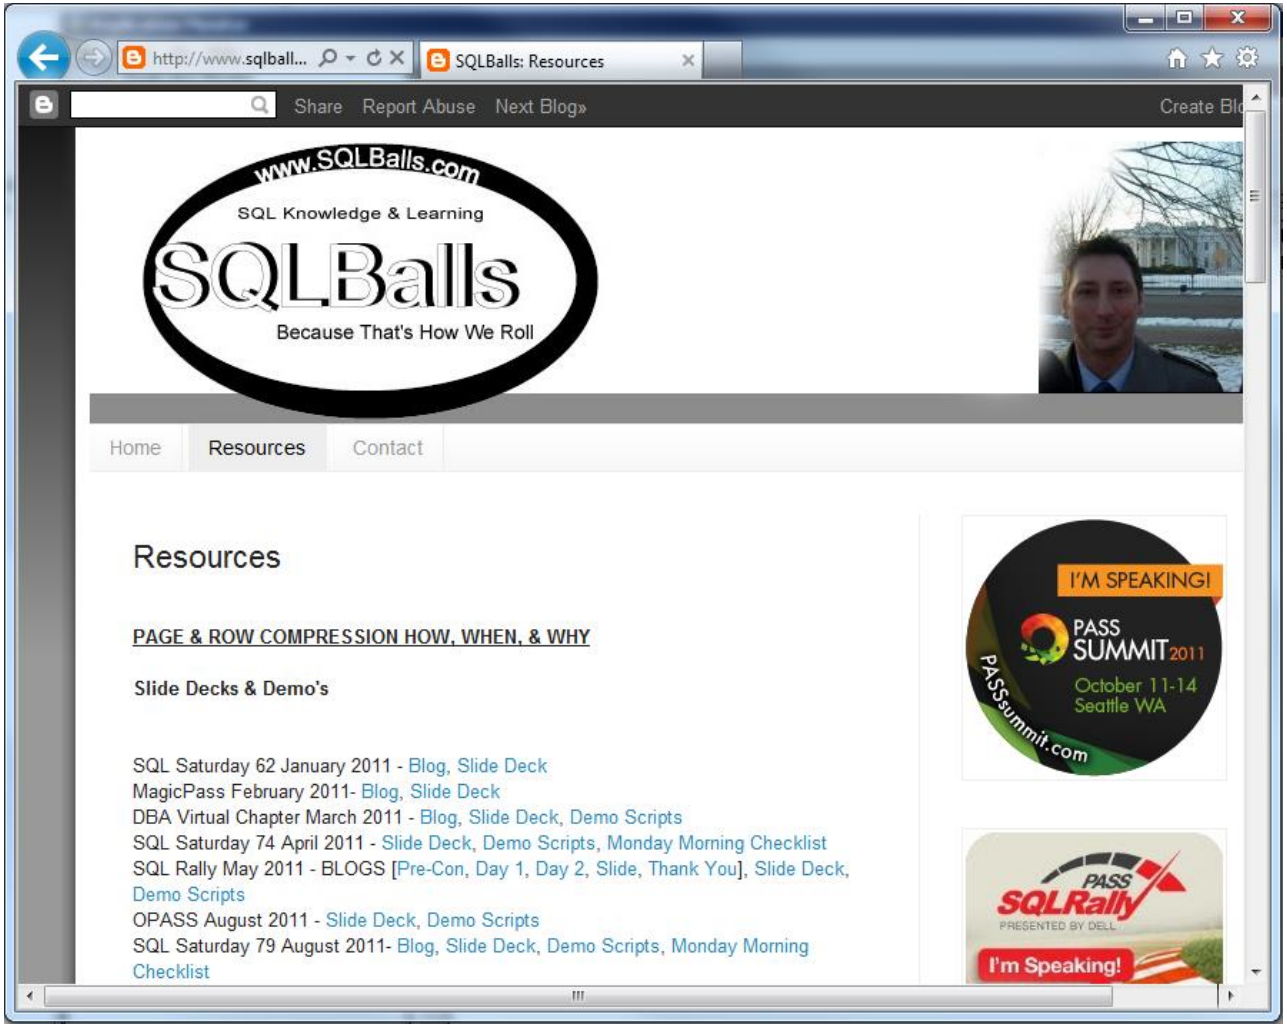
http://www.sqlball...
SQLBalls: Resources
Share
Report Abuse
Next Blog»
www.SQLBalls.com
SQL Knowledge & Learning
SQLBalls
Because That's How We Roll
Home
Resources
Contact
Resources
PAGE & ROW COMPRESSION HOW, WHEN, & WHY
Slide Decks & Demo's
SQL Saturday 62 January 2011 - Blog, Slide Deck
MagicPass February 2011- Blog, Slide Deck
DBA Virtual Chapter March 2011 - Blog, Slide Deck, Demo Scripts
SQL Saturday 74 April 2011 - Slide Deck, Demo Scripts, Monday Morning Checklist
SQL Rally May 2011 - BLOGS [Pre-Con, Day 1, Day 2, Slide, Thank You], Slide Deck, Demo Scripts
OPASS August 2011 - Slide Deck, Demo Scripts
SQL Saturday 79 August 2011- Blog, Slide Deck, Demo Scripts, Monday Morning Checklist
I'M SPEAKING!
PASS SUMMIT 2011
October 11-14
Seattle WA
PASSsummit.com
PASS SQLRally
PRESENTED BY DELL
I'm Speaking!

Questions?