

SO...

- You get that call
- You know the one
- At 3:30 am
- “I’m glad you are awake”
- “I think there is a problem with ...replication”
- “Oh yeah, everything needs to be working by 6:00 am for the data warehouse load”
- “They said they made some changes to things yesterday. Not sure what things but things happened and things aren’t happening now”
- “Thanks, let us know every 10 or so minutes on how things are going”
- “Good Luck”

So what do you do now!!!



Help! My Replication Has Fallen and We Can't Get it Replicating again

Thomas M. Lane, Ph.D.

doclane@gmail.com

@doclane

www.Solve100.com

Thomas M. Lane, PhD.
Husband, Father, Step-Father, Grandfather
BI Consultant

Ph.D. in Chemistry

- Synthetic/Physical Inorganic Chemistry
- Programming (Fortran, Pascal, Basic, C and C++)
- Instrumentation and Quantum Calculation programming
- Helped Automated Four Circle X-ray Diffractometer

Marquette University

- Synthetic/Physical Organosilicone Chemistry
- 3 Years of ...
- **Continuing Education Instructor**

Private Teaching Firm

- I would teach anything with one week or less to prepare
- Everything from Word to C++
- One year

Developer : *Tushaus Computer Services*

- Visual Basic
- Access
- SQL Server
- InstallShield

DBA

- VLDB
- Data Warehousing
- ETL
- Replication
- +12 years
- 7. 2000, 2005, 2008, 2008 R2, 2012, 2014

Data Warehouse

- Data Architect
- ETL
- SSRS, Power
- 15 years
- 2005, 2008 R2, 2012, 2014, 2016

Before we Begin

- Only an Hour Talk so I am only to look at Transactional Replication
- In order to trouble shoot you need to know all the parts
- Lots and lots of moving parts
 - All working together
 - To make one product with many flavors
- Data about the publications is stored in tables in distribution database
- Data managed in the transaction logs and MSrepl_commands



Parts of Replication

Follows a magazine/book producing terminology

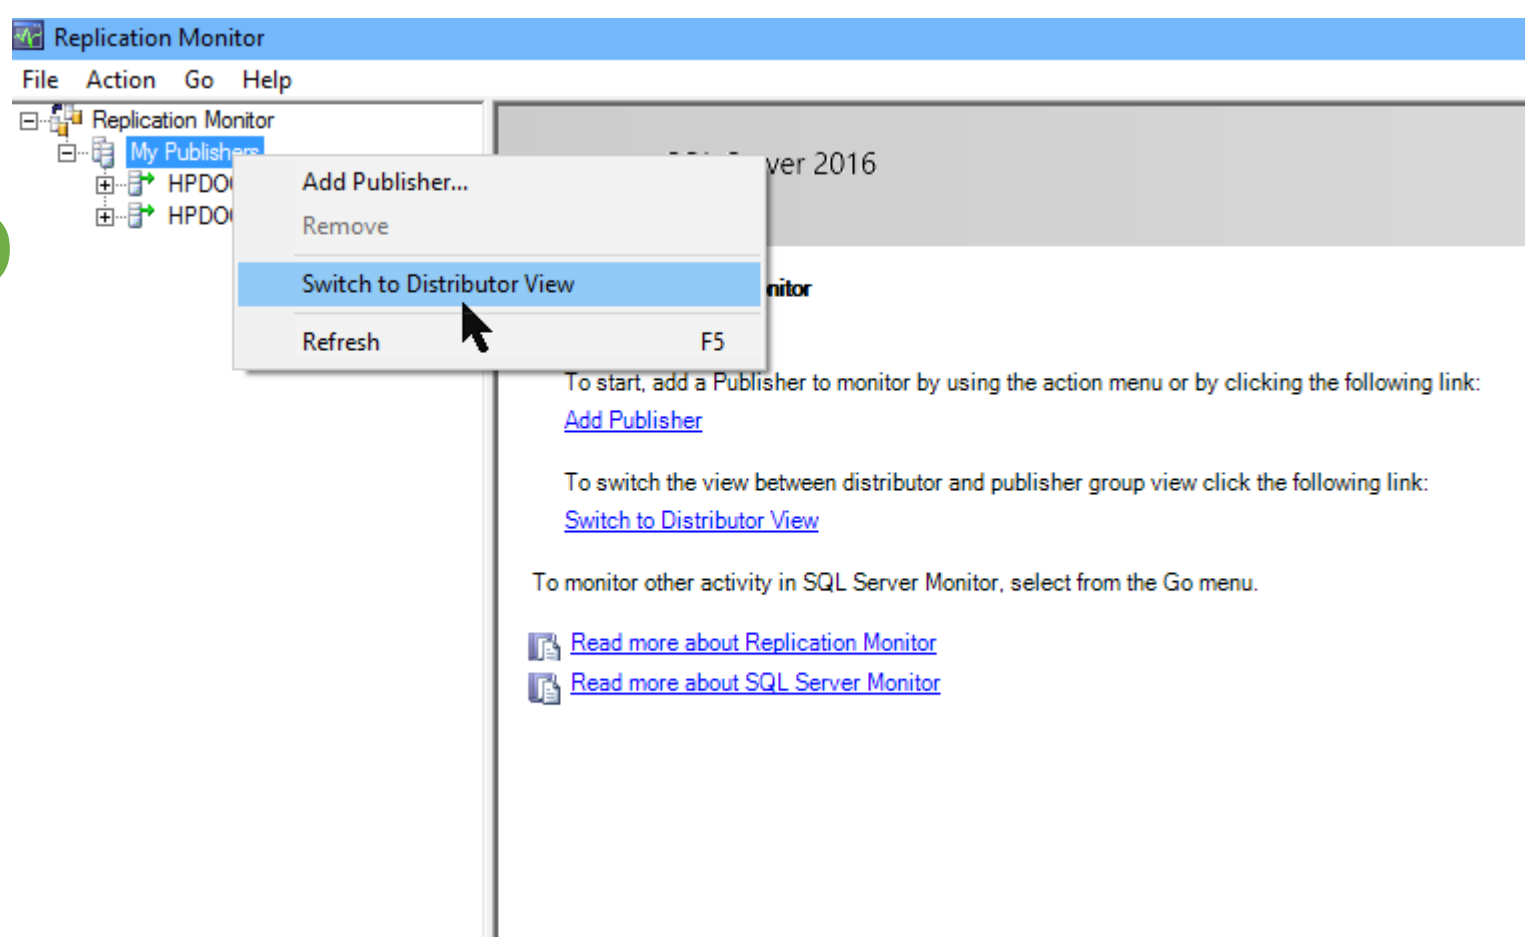
- **Articles**
 - Each article corresponds to a single SQL Server object, or a subset of an object. The objects that are replicated most often are tables, views and stored procedures.
- **The Publication**
 - A group of articles that logically belong together can be combined into a publication. The publication has options defined that apply to all the articles in that publication. The main option defined by a publication is the type of replication that is to be used.
- **The Publisher**
 - The publisher monitors all articles for changes, and makes information about those changes available to the distributor

Parts of Replication

- **The Distributor**
 - SQL Server instance keeps track of all subscribers and all published changes and makes sure that each subscriber gets notified of each change. Most of the changes are tracked in a distribution database. The distributor can be a separate SQL Server instance, but often the distribution service runs on the same machine as the publisher.
- **Subscription**
 - A subscription defines which server (subscriber) is to receive the updates published in a publication.
 - Link between one publication and one subscriber.
 - Push subscriptions and pull subscriptions.
- **Subscriber**
 - The subscriber is the SQL Server instance that receives all the published information through subscriptions.

Lets fix that pro

- Where do we begin
 - Calm your nerves
 - A shot of Tequila
 - Or Two
 - Check LinkedIn
 - Open SSMS
 - Connect to your publisher
 - You can change view
 - Find bad publication
 - If you don't see an error check the subscribers
 - View details
 - Click on the most recent error message



Lets see what is wrong with the Subscription

- Get the transaction sequence number from error message
 - 0x0000004900000388000400000000
- Run `FindProblemReplicationCode.sql`
- Latency
 - Tokens vs Canary Table
- Subscriber data changes

Why did you let this happen!!!

- Why would someone call YOU to tell you there is an issue with replication?
 - This should never happen!
 - Get notification
 - Tracer Token
 - Canary
 - Add a table to publication
 - Only one record
 - Update record every
 - Which is best?
 - In transactional replication subscriber is updatable.
 - Can you make it read-only?
 - From some?



Canary

- Add a table named `dbo.Canary_PubName` to each publication
- `dbo.Canary_PubName` has a single row with a datetime column in it
- A SQL Server Agent job on the publisher updates the datetime to the current timestamp every minute
- A SQL Server Agent job on the subscriber checks `dbo.Canary_PubName` every minute and alerts if the difference between the current time and the timestamp is greater than N minutes



<https://www.brentozar.com/archive/2014/07/monitoring-sql-server-transactional-replication/>

Why did you let this happen!!!

- Run `GenerateTriggers.sql`
 - Leave access for replication user and a sql login
 - Password and user account should be hard to remember and perhaps warn the user.
 - Make sure they know what they are doing



Best Practices

- Script out everything and keep good copies
 - Make sure to add passwords N'password' in script
 - Test them
- Make replication user
 - Minimum required limited access to everything in replication
 - Password does not change
 - User can't change password
 - Domain account
 - May require cred and proxy setup
- Monitor
 - Latency
 - Errors

Best Practices

- Limit changes to publications
 - These will not require re-initialization of snapshot (invasive)
 - ALTER TABLE
 - ALTER VIEW
 - ALTER PROCEDURE
 - ALTER FUNCTION
 - ALTER TRIGGER
 - Consider leaving old column by adding defaults
- Limit the number of row and columns replicated
- Remember each process run an application that uses resources
 - Important when considering Pull or Push subscriptions
- Indexes and Keys are not replicated
 - Use on either side to improve performance

Best Practices

- Consider optimizing the Agent Profiles
- Name all table constraints

General Considerations

- Could not find stored procedure 'sp_MSins_<table_name>' run:
 - sp_scriptPublicationcustomprocs 'name of pub'
- Explicitly adding, dropping, or altering indexes is not supported. Indexes created implicitly for constraints (such as a primary key constraint) are supported.
- Altering or dropping identity columns that are managed by replication is not supported. It is recommended that constraints be explicitly named.
- If a table is published for replication, it is not possible to alter a column in that table to a data type of XML if a publication snapshot has already been generated. To alter the column, you must first remove replication.

General Considerations

- Schema changes should be made only at the Publisher
- Changes propagated to a republishing Subscriber are by default propagated to its Subscribers.
- If the schema change references objects or constraints existing on the Publisher but not on the Subscriber, the schema change will succeed on the Publisher but will fail on the Subscriber.
- All objects on the Subscriber that are referenced when adding a foreign key must have the same name and owner as the corresponding object on the Publisher.

Lets create a Distributor

- Best practice is to have a separate distributor from publisher and subscribers
 - 3 nodes minimum
- Use Wizard but script and save file
 - Always change passwords in replication script files and resave
 - Null replace with N'*password*'
- Check Job list before
- Right click replication (HPDOCLAPTOP\SUBSCRIBER)
 - This snapshot folder does not support **pull** subscriptions created at the Subscriber. It is not a network path or it is a drive letter mapped to a network path. To support both push and pull subscriptions, use a network path to refer to this folder.
- Now look at the Jobs
- 6 New Jobs

Thank you for your time

- Thomas Lane
 - @doclane
 - doclane@gmail.com
- Police: Canary in a Coal Mine
- Monty Python: Monty Python and the Holy Grail

Lets create a Distributor

- Agent history clean up: distribution
 - EXEC dbo.sp_MShistory_cleanup @history_retention = 48
 - By default the SQL Replication Agent History cleanup job runs every 10 minutes and deletes any historical data older the 48 hours.
 - Change? Replication runs every 2 hours, connections bad, etc.?
- Distribution clean up: distribution
 - EXEC dbo.sp_MSdistribution_cleanup @min_distretention = 0, @max_distretention = 72
 - The purpose of this code is to clean up any commands in the MSRepl_Commands table which are more than 72 hours old.
 - NOTE FOR PUBLICATION: If the immediate_sync column is set to 1 then all commands are being held for 72 hours. However, if it is set to 0 then commands are removed from MSRepl_commands the moment that they have been passed to all subscribers... not held for the full 72 hours.
 - exec dbo.sp_helppublication @publication = ''
 - sp_changepublication @publication = 'Test', @property = 'immediate_sync', @value = 'false'

Lets create a Distributor

- Agent history clean up: distribution
 - EXEC sys.sp_expired_subscription_cleanup
 - Checks the status of all the subscriptions of every publication and drops those that have expired.
 - Change?
- Reinitialize subscriptions having data validation failures
 - exec sys.sp_MSreinit_failed_subscriptions @failure_level = 1
 - This agent will check that the subscription is valid
- Replication agents checkup
 - sys.sp_replication_agent_checkup @heartbeat_interval = 10
 - this agent will identify agents that are not logging. It runs every 10 minutes.
- Replication monitoring refresher for distribution
 - exec dbo.sp_replmonitorrefreshjob
 - This agent supports the replication monitor.

Lets create a Distributor

- Creates distribution database
 - [dbo].[MSpublications]
 - Each publication
 - [dbo].[MSarticles]
 - Each Article
 - [dbo].[MSsubscriptions]
 - Each article in each subscription

Lets create a Publication

- New Publication
 - Setup a new Domain account with access to the snapshot location
Login on Publication and distributor SQL Server
USE [master]
GO
CREATE LOGIN [HPDOCLAPTOP\ReplicationUser] FROM WINDOWS WITH DEFAULT_DATABASE=[master]
GO
USE [distribution]
GO
CREATE USER [HPDOCLAPTOP\ReplicationUser] FOR LOGIN [HPDOCLAPTOP\ReplicationUser]
GO
USE [distribution]
GO
ALTER ROLE [db_owner] ADD MEMBER [HPDOCLAPTOP\ReplicationUser]
GO
 - May need to create a credential
 - CREATE CREDENTIAL [HPDOCLAPTOP\ReplicationUser] WITH IDENTITY = N'HPDOCLAPTOP\ReplicationUser', SECRET = N'D1gestiveJuices'
 - Setup SQL Server Agent Replication proxy

Lets create a Publication

- Two new Jobs created
 - Runs separate application not part of SQL Server Engine
 - Logreader Agent
 - {Servername}\{Instance}-Databasename-{1}
 - LogRead.exe
 - Extracts relevant committed transactions from the publisher's log, repackages them and sends them to the distributor in the correct sequence.
 - Publisher [HPDOCLAPTOP\SQL16] -PublisherDB [ReplTransDB] -Distributor [HPDOCLAPTOP\SQL16] -DistributorSecurityMode 1 -Continuous
 - Snapshot Agent
 - {Servername}\{Instance}-Databasename-{Pubname}-{1}
 - -Publisher [HPDOCLAPTOP\SQL16] -PublisherDB [ReplTransDB] -Distributor [HPDOCLAPTOP\SQL16] -Publication [SQL16_ReplTransPub] -DistributorSecurityMode 1
 - Check out Replication Monitor
 - Check out Task manager

Lets create a Subscription

- Create database on Subscriber server
- Add replication user
- Right click on the publication, click New Subscription
- If issues exists
 - Confirm replication user password
 - Make sure to change the package owner to SA
 - Change the “run agent” step to execute as the proxy account
 - Make sure replication user has access to snapshot folder (Full rights)
 - Make sure replication user has access to C:\Program Files\Microsoft SQL Server\1x0\COM (exec)
- Distribution Agent
 - Distrib.exe takes the snapshots, and log entries from the agents we’ve described, and dispatches them to the subscribers.
 - Check Task Manager

Lets create a Subscription

- Snapshot needs to be run at least once since the oldest transaction aged out
- If distribution database gets large change the two jobs to clean up more often
- Look are replication Monitor
- Look at agent status
- Tracers – can be done programmatically
- Question: How good is this tracer? Can you depend on it?
- Check out Warning
 - Configure alerts
 - Will need operator set up

Lets see what is wrong with the Subscription

- Run `GenerateTriggers.sql`
 - A access hole for one sql server login
 - Make sure they know what they are doing

debugging

```
sp_helptext 'sp_MSdel_dboF4201'
```

```
{CALL [sp_MSdel_dboF4201] (789211)}
```

```
CALL [sp_MSdel_dboF4201] (778870)
```

```
SET IDENTITY_INSERT [dbo].[F4211] ON
```

```
INSERT INTO [dbo].[F4211]
```

```
([SDKC00]
```

```
, [SDDOCO]
```

```
, [SDDCTO]
```

```
, [SDLNID]
```

```
, [F4211_ID]
```

```
)
```

```
VALUES('xxx', 0, 'xx', 0, 789211)
```

```
SET IDENTITY_INSERT [dbo].[F4211] OFF
```

```
select * from [dbo].[F4211] where [F4211_ID] = 789211
```

```
select *
```

```
from msrepl_commands (nolock)
```

```
where command_id = <no in message>
```